

Customer Churn Prediction Analysis Using Supervised Learning Approach

- Using Normalized Data
- RFE Feature Selection Technique
- Using SMOTE Technique to handle class imbalance
- Models: Logistic Regression, Random Forest and SVM Model
- Using ROC AUC Curve to evaluate the model's performance.

Define the Problem:

First we will define the objectives of this analysis and the questions you want to answer using the data and understand the context and purpose of the analysis.

- The analysis is to develop predictive model to forecast churn in telecommunication businesses

Importing Necessary Liabries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
from sklearn.preprocessing import LabelEncoder
from sklearn.feature_selection import RFE
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.feature_selection import SelectKBest, chi2
from sklearn import metrics
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, C
import matplotlib.pyplot as plt
from imblearn.over_sampling import SMOTE
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import roc_curve, roc_auc_score
```


Importing Dataset

```
In [2]: data = pd.read_csv('Customer Churn Dataset.csv')
```

```
In [3]: data.head()
```

Out[3]:	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	Inte
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	
1	5575-GNVDE	Male	0	No	No	34	Yes	No	
2	3668-QPYBK	Male	0	No	No	2	Yes	No	
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	
4	9237-HQITU	Female	0	No	No	2	Yes	No	

5 rows × 21 columns



Data Cleaning and Preprocessing

In this section we will;

- Handle missing values: Identify and deal with missing data by imputation or removal.
- Remove duplicates (if any): Eliminate duplicate records if present in the dataset (if any)
- Standardize data formats: Ensure consistency in data formats and units.
- Feature engineering: Create new features or transform existing ones to better represent the data and improve model performance.

In [4]: `data.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                 7043 non-null   object
2   SeniorCitizen          7043 non-null   int64
3   Partner                7043 non-null   object
4   Dependents             7043 non-null   object
5   tenure                 7043 non-null   int64
6   PhoneService           7043 non-null   object
7   MultipleLines          7043 non-null   object
8   InternetService        7043 non-null   object
9   OnlineSecurity         7043 non-null   object
10  OnlineBackup           7043 non-null   object
11  DeviceProtection       7043 non-null   object
12  TechSupport            7043 non-null   object
13  StreamingTV            7043 non-null   object
14  StreamingMovies        7043 non-null   object
15  Contract               7043 non-null   object
16  PaperlessBilling       7043 non-null   object
17  PaymentMethod          7043 non-null   object
18  MonthlyCharges         7043 non-null   float64
19  TotalCharges           7043 non-null   object
20  Churn                  7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB

```

```
In [5]: data.size
```

```
Out[5]: 147903
```

```
In [6]: data.nunique()
```

```

Out[6]: customerID            7043
gender                 2
SeniorCitizen          2
Partner                2
Dependents             2
tenure                 73
PhoneService           2
MultipleLines          3
InternetService        3
OnlineSecurity         3
OnlineBackup           3
DeviceProtection       3
TechSupport            3
StreamingTV            3
StreamingMovies        3
Contract               3
PaperlessBilling       2
PaymentMethod          4
MonthlyCharges         1585
TotalCharges           6531
Churn                  2
dtype: int64

```

```
In [7]: data.dtypes
```

```
Out[7]: customerID      object
gender      object
SeniorCitizen  int64
Partner      object
Dependents    object
tenure      int64
PhoneService  object
MultipleLines object
InternetService object
OnlineSecurity object
OnlineBackup  object
DeviceProtection object
TechSupport  object
StreamingTV   object
StreamingMovies object
Contract      object
PaperlessBilling object
PaymentMethod object
MonthlyCharges float64
TotalCharges  object
Churn         object
dtype: object
```

Handling Missing Data

```
In [8]: data.isnull().sum()
```

```
Out[8]: customerID      0
gender      0
SeniorCitizen  0
Partner      0
Dependents    0
tenure      0
PhoneService  0
MultipleLines  0
InternetService  0
OnlineSecurity  0
OnlineBackup  0
DeviceProtection  0
TechSupport  0
StreamingTV   0
StreamingMovies  0
Contract      0
PaperlessBilling  0
PaymentMethod  0
MonthlyCharges  0
TotalCharges  0
Churn         0
dtype: int64
```

We noticed that Total Charges is represented as a categorical variable instead of numeric, the code below converts it to numeric variable

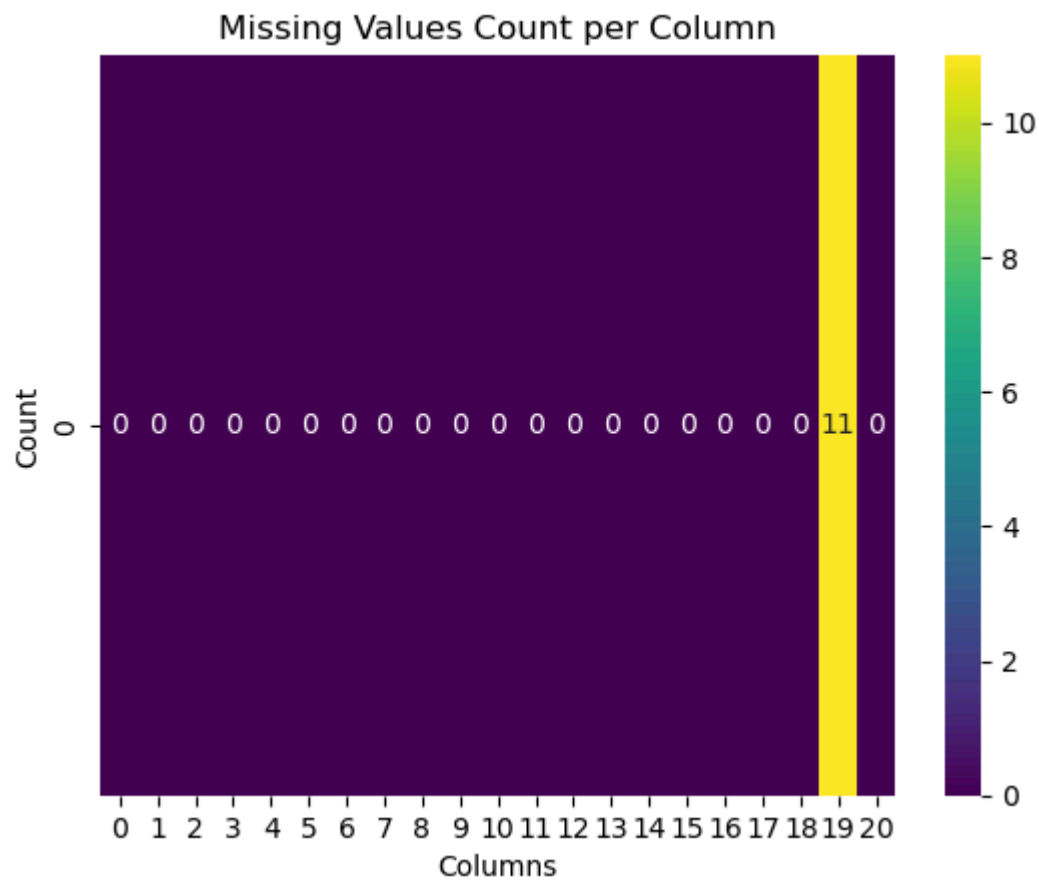
```
In [9]: data.TotalCharges = pd.to_numeric(data.TotalCharges, errors='coerce')
```

```
In [10]: data.isnull().sum()
```

```
Out[10]: customerID      0
gender      0
SeniorCitizen  0
Partner      0
Dependents   0
tenure       0
PhoneService  0
MultipleLines  0
InternetService  0
OnlineSecurity  0
OnlineBackup  0
DeviceProtection  0
TechSupport   0
StreamingTV   0
StreamingMovies  0
Contract      0
PaperlessBilling  0
PaymentMethod  0
MonthlyCharges  0
TotalCharges  11
Churn         0
dtype: int64
```

```
In [11]: # Calculate the count of missing values in each column and convert it to a 2D array
missing_values_count = data.isna().sum().values.reshape(1, -1)

# Create a heatmap to visualize missing values count
sns.heatmap(missing_values_count, annot=True, cmap='viridis')
plt.title('Missing Values Count per Column')
plt.xlabel('Columns')
plt.ylabel('Count')
plt.show()
```



Total Charges have 11 missing data. We will drop the enter columns

```
In [12]: #Dropping the missing values in Total Charges column
data.dropna(subset=['TotalCharges'], inplace=True)
```

```
In [13]: data.isna().sum()
```

```
Out[13]: customerID      0
gender      0
SeniorCitizen  0
Partner     0
Dependents  0
tenure      0
PhoneService  0
MultipleLines  0
InternetService  0
OnlineSecurity  0
OnlineBackup  0
DeviceProtection  0
TechSupport  0
StreamingTV  0
StreamingMovies  0
Contract     0
PaperlessBilling  0
PaymentMethod  0
MonthlyCharges  0
TotalCharges  0
Churn        0
dtype: int64
```

```
In [14]: data.drop(['customerID'], axis=1, inplace=True)
```

```
In [15]: data.head()
```

```
Out[15]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService
0	Female	0	Yes	No	1	No	No phone service	DSL
1	Male	0	No	No	34	Yes	No	DSL
2	Male	0	No	No	2	Yes	No	DSL
3	Male	0	No	No	45	No	No phone service	DSL
4	Female	0	No	No	2	Yes	No	Fiber optic

Checking for Duplicates

```
In [16]: duplicate_rows = data[data.duplicated()]

if len(duplicate_rows) == 0:
    print("No duplicate rows found.")
else:
    print("Duplicate Rows:")
    print(duplicate_rows)
```

Duplicate Rows:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	\
964	Male	0	No	No	1	Yes	
1338	Male	0	No	No	1	Yes	
1491	Female	0	No	No	1	Yes	
1739	Male	0	No	No	1	Yes	
1932	Male	0	No	No	1	Yes	
2713	Male	0	No	No	1	Yes	
2892	Male	0	No	No	1	Yes	
3301	Female	1	No	No	1	Yes	
3754	Male	0	No	No	1	Yes	
4098	Male	0	No	No	1	Yes	
4476	Female	0	No	No	1	Yes	
5506	Male	0	No	No	1	Yes	
5736	Male	0	No	No	1	Yes	
5759	Female	0	No	No	1	Yes	
6267	Female	0	No	No	1	Yes	
6499	Male	0	No	No	1	Yes	
6518	Male	0	No	No	1	Yes	
6609	Male	0	No	No	1	Yes	
6706	Female	0	No	No	1	Yes	
6764	Female	0	No	No	1	Yes	
6774	Female	0	No	No	1	Yes	
6924	Male	0	No	No	1	Yes	

	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	\
964	No	DSL	No	No	
1338	No	No	No internet service	No internet service	
1491	No	No	No internet service	No internet service	
1739	No	Fiber optic	No	No	
1932	No	No	No internet service	No internet service	
2713	No	No	No internet service	No internet service	
2892	No	No	No internet service	No internet service	
3301	No	Fiber optic	No	No	
3754	No	No	No internet service	No internet service	
4098	No	No	No internet service	No internet service	
4476	No	No	No internet service	No internet service	
5506	No	No	No internet service	No internet service	
5736	No	No	No internet service	No internet service	
5759	No	Fiber optic	No	No	
6267	No	Fiber optic	No	No	
6499	No	No	No internet service	No internet service	
6518	No	DSL	No	No	
6609	No	No	No internet service	No internet service	
6706	No	No	No internet service	No internet service	
6764	No	Fiber optic	No	No	
6774	No	No	No internet service	No internet service	
6924	No	Fiber optic	No	No	

	DeviceProtection	TechSupport	StreamingTV	\
964	No	No	No	
1338	No internet service	No internet service	No internet service	
1491	No internet service	No internet service	No internet service	
1739	No	No	No	
1932	No internet service	No internet service	No internet service	
2713	No internet service	No internet service	No internet service	
2892	No internet service	No internet service	No internet service	
3301	No	No	No	
3754	No internet service	No internet service	No internet service	
4098	No internet service	No internet service	No internet service	

4476	No internet service	No internet service	No internet service
5506	No internet service	No internet service	No internet service
5736	No internet service	No internet service	No internet service
5759	No	No	No
6267	No	No	No
6499	No internet service	No internet service	No internet service
6518	No	No	No
6609	No internet service	No internet service	No internet service
6706	No internet service	No internet service	No internet service
6764	No	No	No
6774	No internet service	No internet service	No internet service
6924	No	No	No

	StreamingMovies	Contract	PaperlessBilling	PaymentMethod \
964	No	Month-to-month	Yes	Mailed check
1338	No internet service	Month-to-month	No	Mailed check
1491	No internet service	Month-to-month	No	Mailed check
1739	No	Month-to-month	Yes	Electronic check
1932	No internet service	Month-to-month	No	Mailed check
2713	No internet service	Month-to-month	Yes	Mailed check
2892	No internet service	Month-to-month	No	Mailed check
3301	No	Month-to-month	Yes	Electronic check
3754	No internet service	Month-to-month	No	Mailed check
4098	No internet service	Month-to-month	Yes	Mailed check
4476	No internet service	Month-to-month	No	Mailed check
5506	No internet service	Month-to-month	No	Mailed check
5736	No internet service	Month-to-month	No	Mailed check
5759	No	Month-to-month	Yes	Mailed check
6267	No	Month-to-month	Yes	Electronic check
6499	No internet service	Month-to-month	No	Mailed check
6518	No	Month-to-month	No	Electronic check
6609	No internet service	Month-to-month	Yes	Mailed check
6706	No internet service	Month-to-month	No	Mailed check
6764	No	Month-to-month	Yes	Electronic check
6774	No internet service	Month-to-month	No	Mailed check
6924	No	Month-to-month	Yes	Electronic check

	MonthlyCharges	TotalCharges	Churn
964	45.70	45.70	Yes
1338	20.15	20.15	Yes
1491	19.55	19.55	No
1739	69.90	69.90	Yes
1932	20.20	20.20	No
2713	20.45	20.45	No
2892	20.45	20.45	No
3301	69.60	69.60	Yes
3754	20.05	20.05	No
4098	20.20	20.20	Yes
4476	20.90	20.90	Yes
5506	20.20	20.20	No
5736	20.05	20.05	No
5759	70.15	70.15	Yes
6267	70.10	70.10	Yes
6499	20.30	20.30	No
6518	45.30	45.30	Yes
6609	20.10	20.10	Yes
6706	19.90	19.90	No
6764	69.20	69.20	Yes
6774	19.65	19.65	No
6924	69.35	69.35	Yes

```
In [17]: #Descriptive Analysis
data.describe(include = 'all')
```

```
Out[17]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	Intern
count	7032	7032.000000	7032	7032	7032.000000	7032	7032	
unique	2	NaN	2	2	NaN	2	3	
top	Male	NaN	No	No	NaN	Yes	No	F
freq	3549	NaN	3639	4933	NaN	6352	3385	
mean	NaN	0.162400	NaN	NaN	32.421786	NaN	NaN	
std	NaN	0.368844	NaN	NaN	24.545260	NaN	NaN	
min	NaN	0.000000	NaN	NaN	1.000000	NaN	NaN	
25%	NaN	0.000000	NaN	NaN	9.000000	NaN	NaN	
50%	NaN	0.000000	NaN	NaN	29.000000	NaN	NaN	
75%	NaN	0.000000	NaN	NaN	55.000000	NaN	NaN	
max	NaN	1.000000	NaN	NaN	72.000000	NaN	NaN	

Now let's convert Senior Citizen column from numeric to categorical labels of ("No" and "Yes") and display the first few rows of the dataframe with the updated values in the "SeniorCitizen" column.

```
In [18]: data.SeniorCitizen = data.SeniorCitizen.map({0: "No", 1: "Yes"})
data.head()
```

```
Out[18]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService
0	Female	No	Yes	No	1	No	No phone service	DSL
1	Male	No	No	No	34	Yes	No	DSL
2	Male	No	No	No	2	Yes	No	DSL
3	Male	No	No	No	45	No	No phone service	DSL
4	Female	No	No	No	2	Yes	No	Fiber optic

Exploratory Data Analysis (EDA):

- Here we will summarize and visualize the data using statistical measures, charts, and graphs.

```
In [19]: data.dtypes
```

```
Out[19]: gender           object
SeniorCitizen          object
Partner                object
Dependents             object
tenure                 int64
PhoneService           object
MultipleLines          object
InternetService        object
OnlineSecurity         object
OnlineBackup           object
DeviceProtection       object
TechSupport            object
StreamingTV            object
StreamingMovies        object
Contract              object
PaperlessBilling       object
PaymentMethod          object
MonthlyCharges         float64
TotalCharges           float64
Churn                  object
dtype: object
```

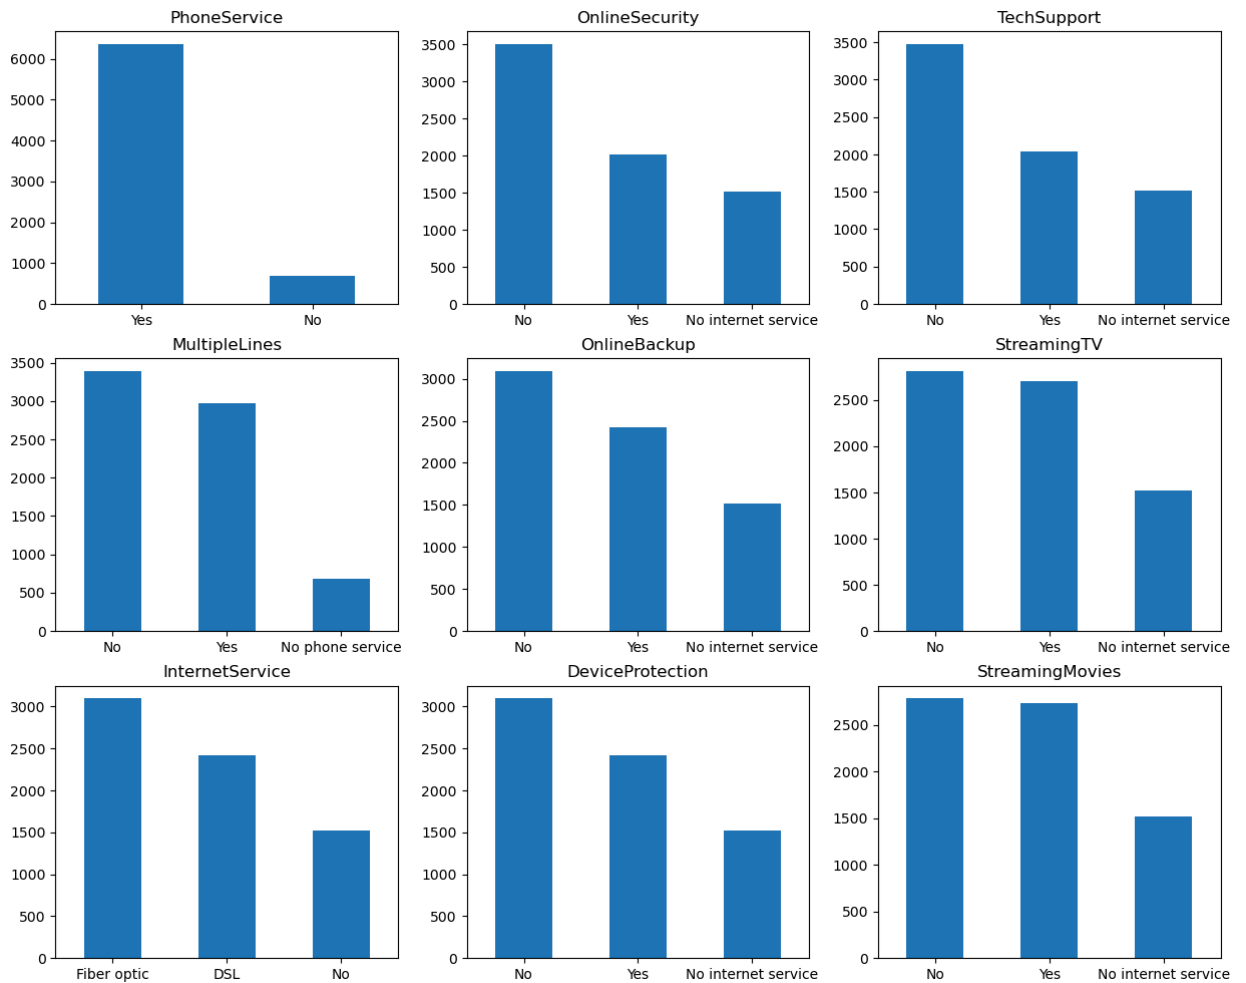
Visualizing the Distribution of Services

```
In [20]: services = ['PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity',
                    'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMov

fig, axes = plt.subplots(nrows = 3, ncols = 3, figsize = (15,12))
for i, item in enumerate(services):
    if i < 3:
        ax = data[item].value_counts().plot(kind = 'bar', ax=axes[i,0], rot = 0)

    elif i >=3 and i < 6:
        ax = data[item].value_counts().plot(kind = 'bar', ax=axes[i-3,1], rot = 0)

    elif i < 9:
        ax = data[item].value_counts().plot(kind = 'bar', ax=axes[i-6,2], rot = 0)
    ax.set_title(item)
```



Using Groupby

In [21]: `data.head()`

Out[21]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService
0	Female	No	Yes	No	1	No	No phone service	DSL
1	Male	No	No	No	34	Yes	No	DSL
2	Male	No	No	No	2	Yes	No	DSL
3	Male	No	No	No	45	No	No phone service	DSL
4	Female	No	No	No	2	Yes	No	Fiber optic

Overall Percentage and Count of Customers that Churned

```
In [22]: #data['Churn'].value_counts()/100
percentage_counts = data['Churn'].value_counts(normalize=True) * 100
print(percentage_counts)
```


```
No      73.421502
Yes     26.578498
Name: Churn, dtype: float64
```

```
In [23]: data['Churn'].value_counts()
```

```
Out[23]: No      5163
Yes     1869
Name: Churn, dtype: int64
```

```
In [24]: # Visualizing and Calculate percentage counts
percentage_counts = data['Churn'].value_counts(normalize=True) * 100

# Plotting
labels = percentage_counts.index
sizes = percentage_counts.values
colors = ['lightskyblue', 'orange']
plt.pie(sizes, labels=labels, colors=colors, autopct='%1.1f%%', startangle=140)
plt.title('Percentage of Churn')
plt.axis('equal')
plt.show()
```




We can see there is a significant class imbalance between customers that churned and customers that didn't churn. In the coming steps we will be using the SMOTE technique an oversampling method to address the class imbalance as this imbalance can affect the model's performance.

Churn by Gender

```
In [25]: # Calculate the count of churned and non-churned customers by gender
churn_count = data.groupby(['gender', 'Churn']).size().unstack()

custom_palette = {"Male": "lightskyblue", "Female": "orange"}

# Plotting
churn_count.plot(kind='bar', stacked=True)
plt.xlabel('Gender')
plt.ylabel('Count')
plt.title('Churn Count by Gender')
plt.legend(title='Churn', loc='upper right')
plt.show()
```



Female and Male have almost the same number of churn and non-churn customers

Churn by Senior Citizen

Customer that are Senior Citizen who Churned

```
In [26]: churn_counts = data[data['Churn'] == 'Yes'].groupby('SeniorCitizen').size()
print(churn_counts)
```

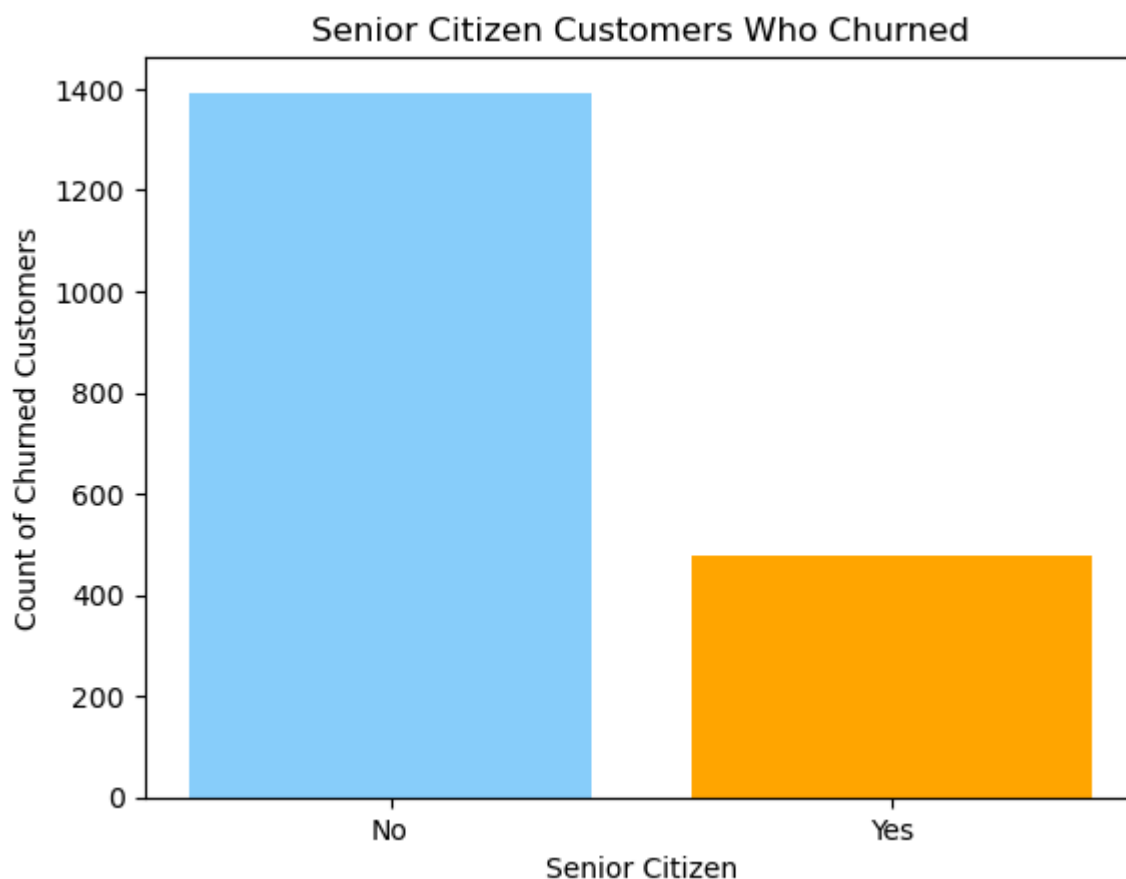
```
SeniorCitizen
No      1393
Yes      476
dtype: int64
```

The above shows that there are 1393 customers who are not senior citizens have churned and also 476 customers who are senior citizens and have churned. In summary, this output tells you how many customers from each group have churned. It provides insight into the churn behavior based on the 'SeniorCitizen' status.

```
In [27]: #Let's visualize it

# Define custom colors
colors = ['lightskyblue', 'orange']

# Plotting
plt.bar(churn_counts.index, churn_counts.values, color=colors)
plt.xlabel('Senior Citizen')
plt.ylabel('Count of Churned Customers')
plt.title('Senior Citizen Customers Who Churned')
plt.xticks(churn_counts.index, ['No', 'Yes']) # Set the x-ticks labels
plt.show()
```



Customer that are Senior Citizen who did not Churn

```
In [28]: churn_counts = data[data['Churn'] == 'No'].groupby('SeniorCitizen').size()
print(churn_counts)
```

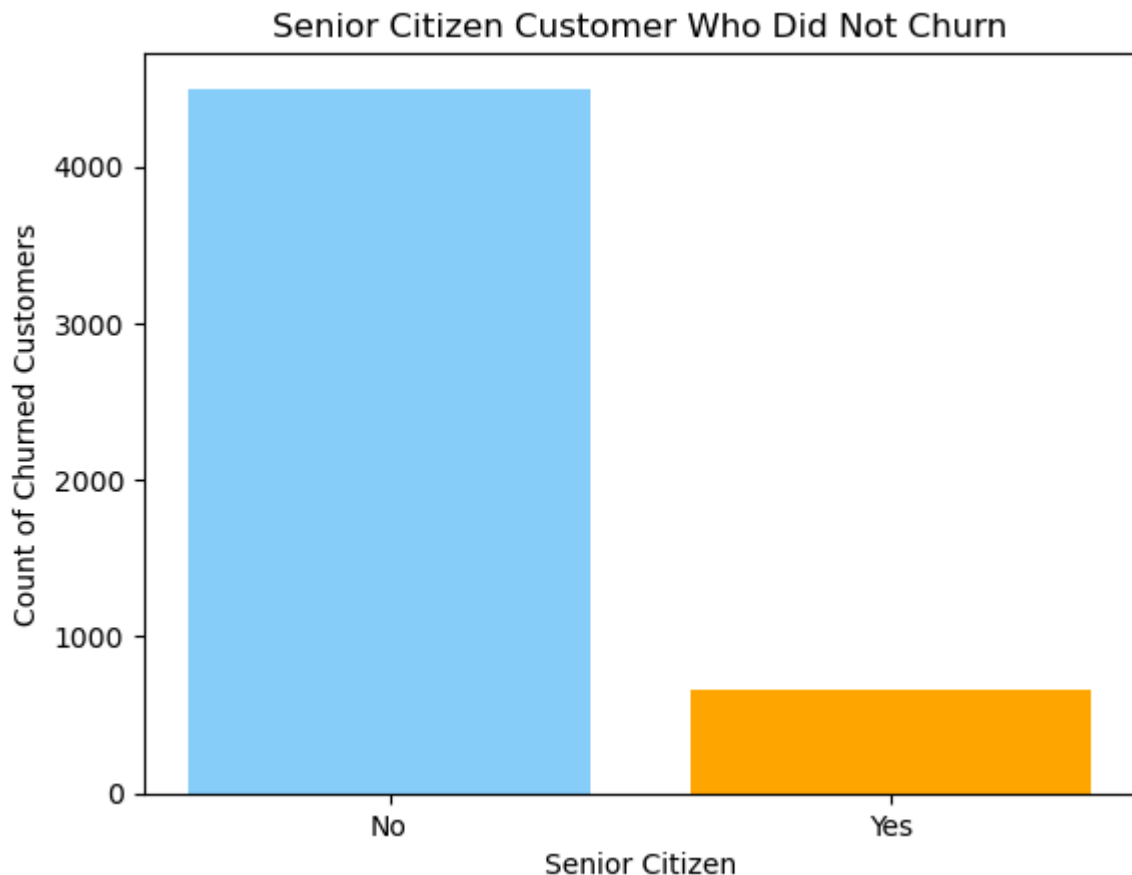
```
SeniorCitizen
No      4497
Yes      666
dtype: int64
```

```
In [29]: #Let's visualize it

# Define custom colors
colors = ['lightskyblue', 'orange']

# Plotting
plt.bar(churn_counts.index, churn_counts.values, color=colors)
plt.xlabel('Senior Citizen')
plt.ylabel('Count of Churned Customers')
plt.title('Senior Citizen Customer Who Did Not Churn')
plt.xticks
```

```
Out[29]: <function matplotlib.pyplot.xticks(ticks=None, labels=None, **kwargs)>
```



Churn by Internet Service Customers

Internet Service Customers who Churned

```
In [30]: churn_count = data[data['Churn'] == 'Yes'].groupby(['InternetService']).size()
print(churn_count)
```



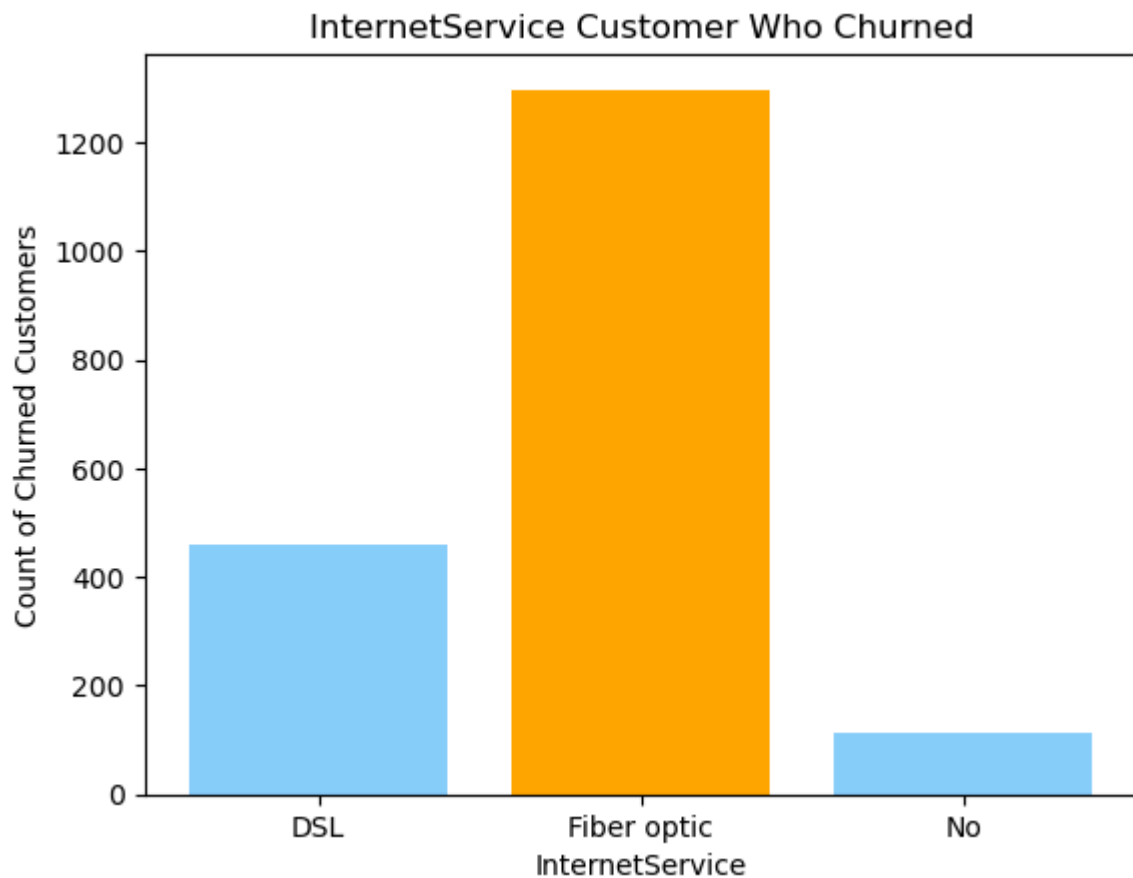
```
InternetService
DSL          459
Fiber optic  1297
No           113
dtype: int64
```

In [31]: *#Let's visualize it*

```
# Define custom colors
colors = ['lightskyblue', 'orange']

# Plotting
plt.bar(churn_count.index, churn_count.values, color=colors)
plt.xlabel('InternetService')
plt.ylabel('Count of Churned Customers')
plt.title('InternetService Customer Who Churned')
plt.xticks
```

Out[31]: <function matplotlib.pyplot.xticks(ticks=None, labels=None, **kwargs)>



Internet Service Customers who did not Churned

In [32]:

```
churn_count = data[data['Churn'] == 'No'].groupby(['InternetService']).size()
print(churn_count)
```

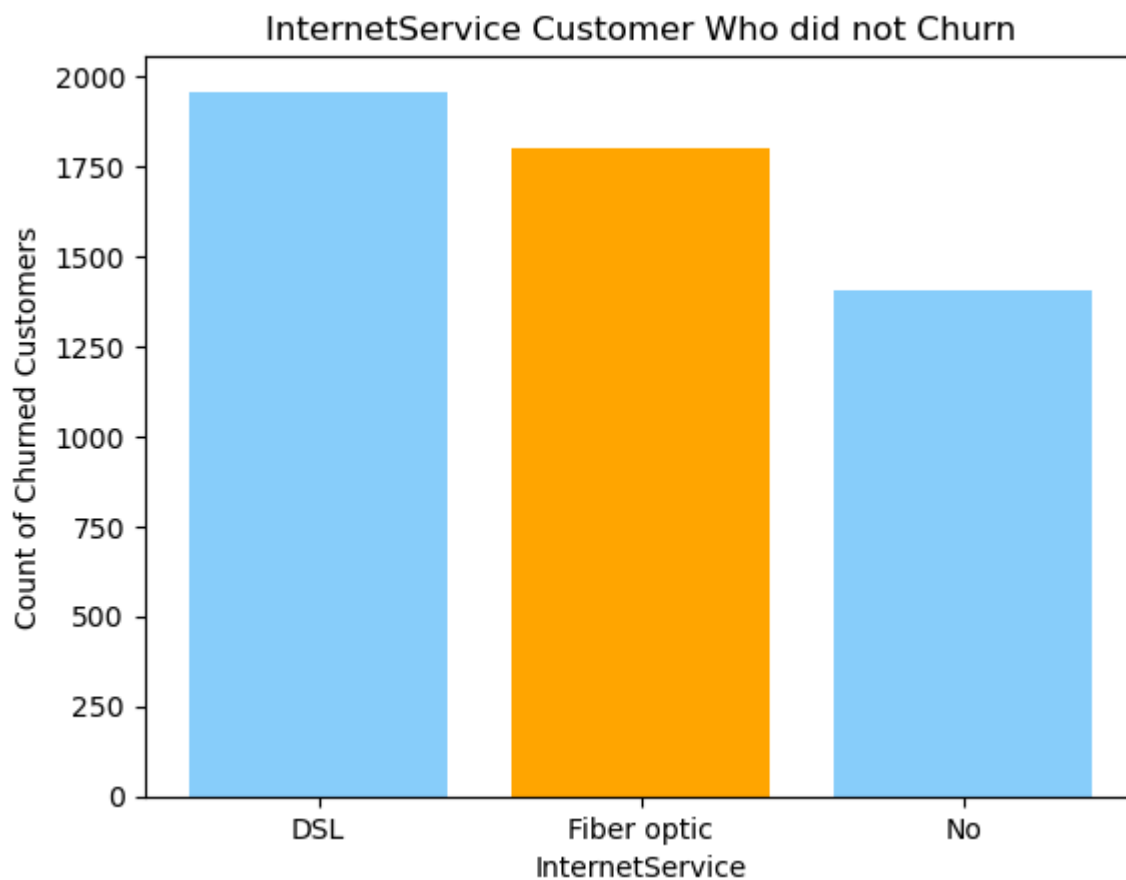
```
InternetService
DSL          1957
Fiber optic  1799
No           1407
dtype: int64
```

In [33]: *#Let's visualize it*

```
# Define custom colors
colors = ['lightskyblue', 'orange']

# Plotting
plt.bar(churn_count.index, churn_count.values, color=colors)
plt.xlabel('InternetService')
plt.ylabel('Count of Churned Customers')
plt.title('InternetService Customer Who did not Churn')
plt.xticks
```

Out[33]: <function matplotlib.pyplot.xticks(ticks=None, labels=None, **kwargs)>



Churn by MultipleLines

MultipleLines Customers who Churned

In [34]: `churn_counts = data[data['Churn'] == 'Yes'].groupby(['MultipleLines']).size()`
`print(churn_counts)`

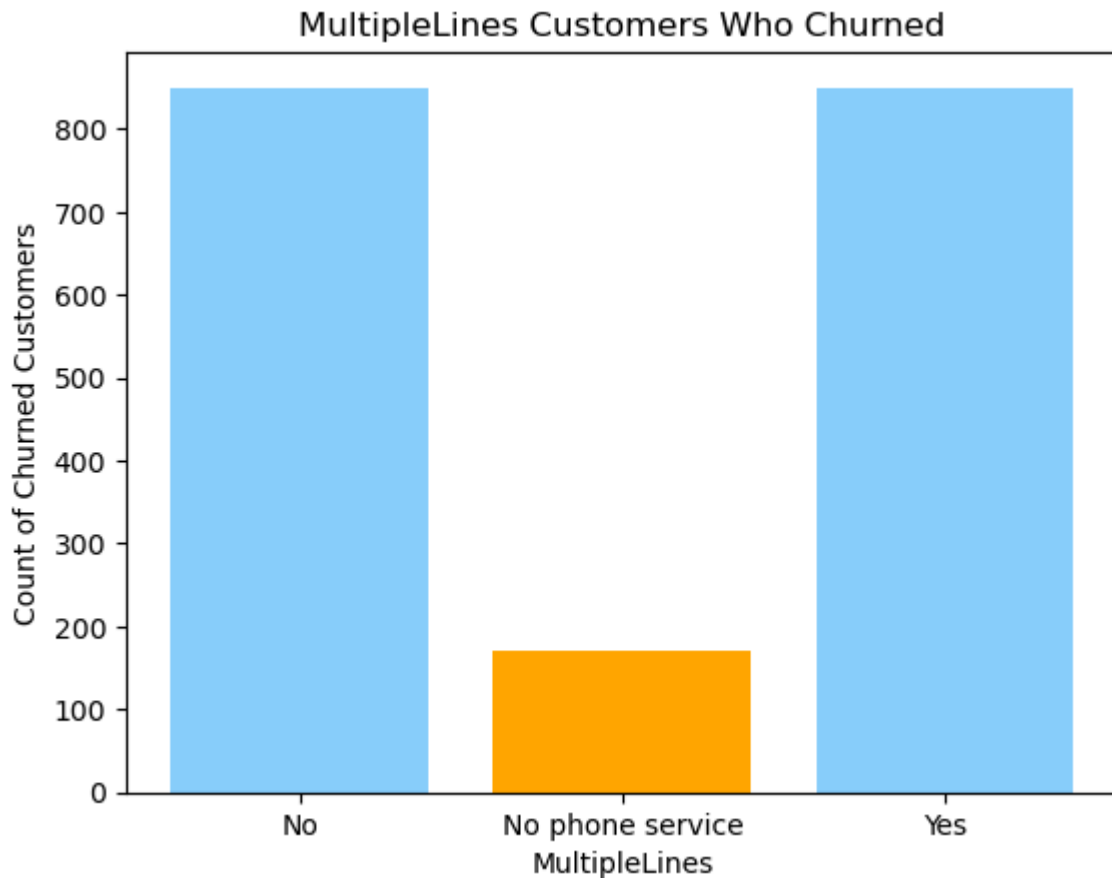
```
MultipleLines
No                849
No phone service   170
Yes               850
dtype: int64
```

In [35]: *#Let's visualize it*

```
# Define custom colors
colors = ['lightskyblue', 'orange']

# Plotting
plt.bar(churn_counts.index, churn_counts.values, color=colors)
plt.xlabel('MultipleLines')
plt.ylabel('Count of Churned Customers')
plt.title('MultipleLines Customers Who Churned')
plt.xticks
```

Out[35]: <function matplotlib.pyplot.xticks(ticks=None, labels=None, **kwargs)>



MultipleLine Customers who did not Churn

In [36]: `churn_counts = data[data['Churn'] == 'No'].groupby(['MultipleLines']).size()`
`print(churn_counts)`

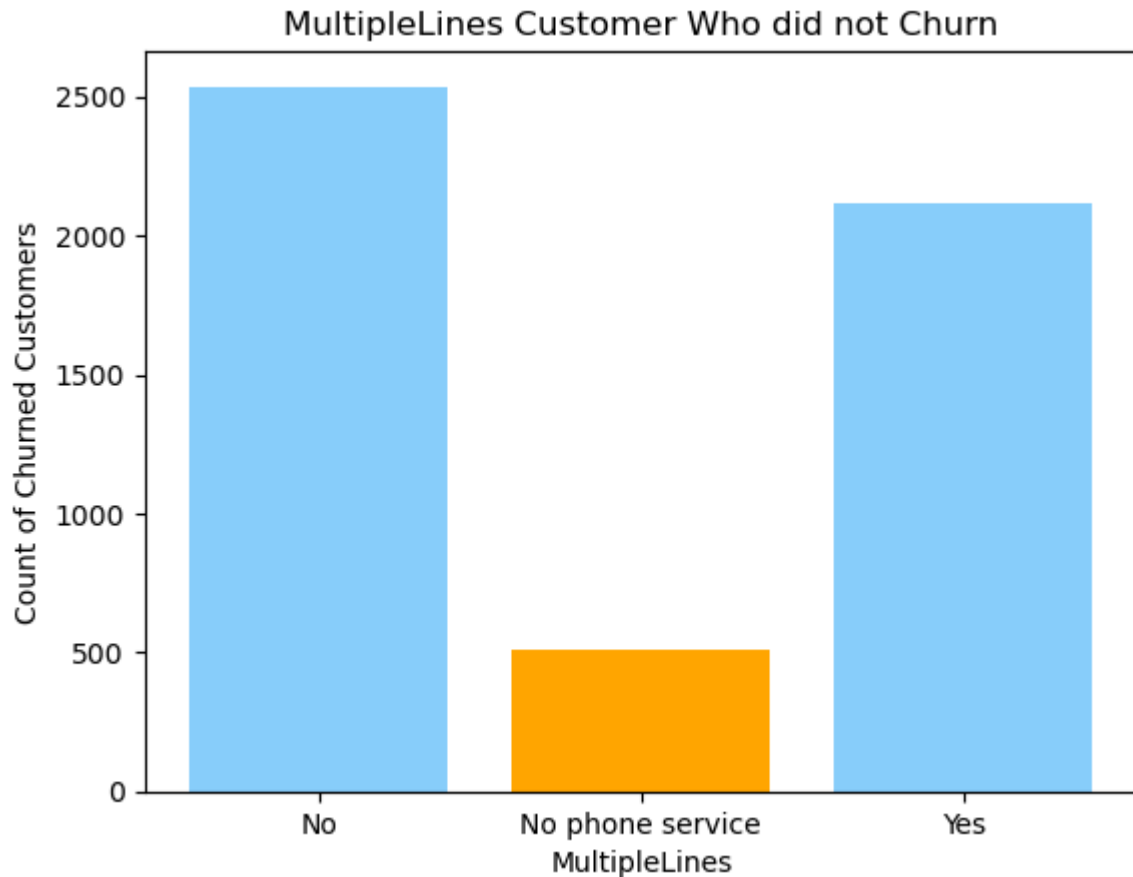
```
MultipleLines
No                2536
No phone service    510
Yes               2117
dtype: int64
```

In [37]: *#Let's visualize it*

```
# Define custom colors
colors = ['lightskyblue', 'orange']
```

```
# Plotting
plt.bar(churn_counts.index, churn_counts.values, color=colors)
plt.xlabel('MultipleLines')
plt.ylabel('Count of Churned Customers')
plt.title('MultipleLines Customer Who did not Churn')
plt.xticks
```

Out[37]: <function matplotlib.pyplot.xticks(ticks=None, labels=None, **kwargs)>



Churn by Contract Type

Contract Type by Customer who Churned

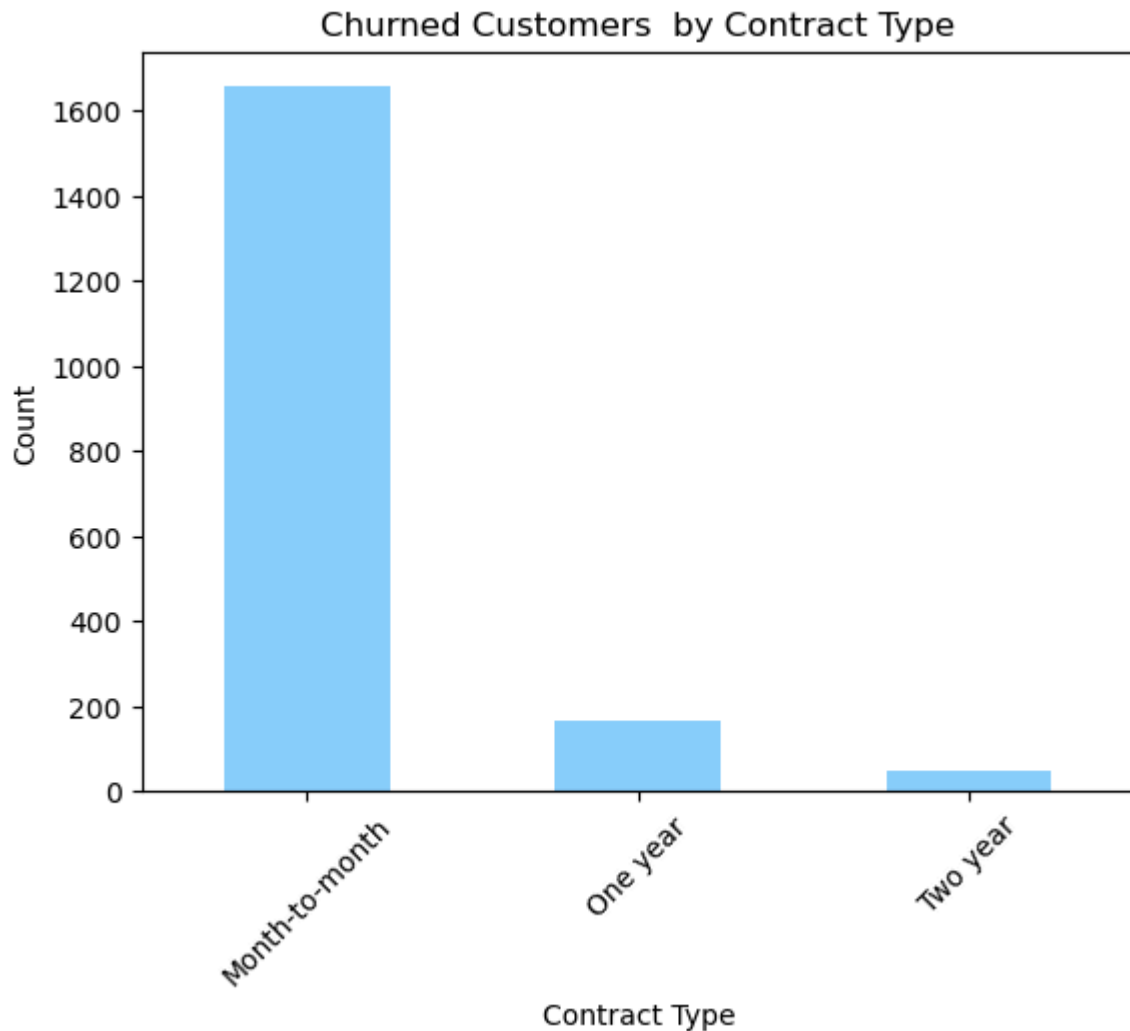
```
In [38]: #Customers who Churned
churn_counts = data[data['Churn'] == 'Yes'].groupby(['Contract']).size()
print(churn_counts)
```

```
Contract
Month-to-month    1655
One year           166
Two year           48
dtype: int64
```

```
In [39]: churn_counts = data[data['Churn'] == 'Yes'].groupby(['Contract']).size()

# Plotting
churn_counts.plot(kind='bar', color='lightskyblue')
plt.xlabel('Contract Type')
```

```
plt.ylabel('Count')
plt.title('Churned Customers by Contract Type')
plt.xticks(rotation=45)
plt.show()
```



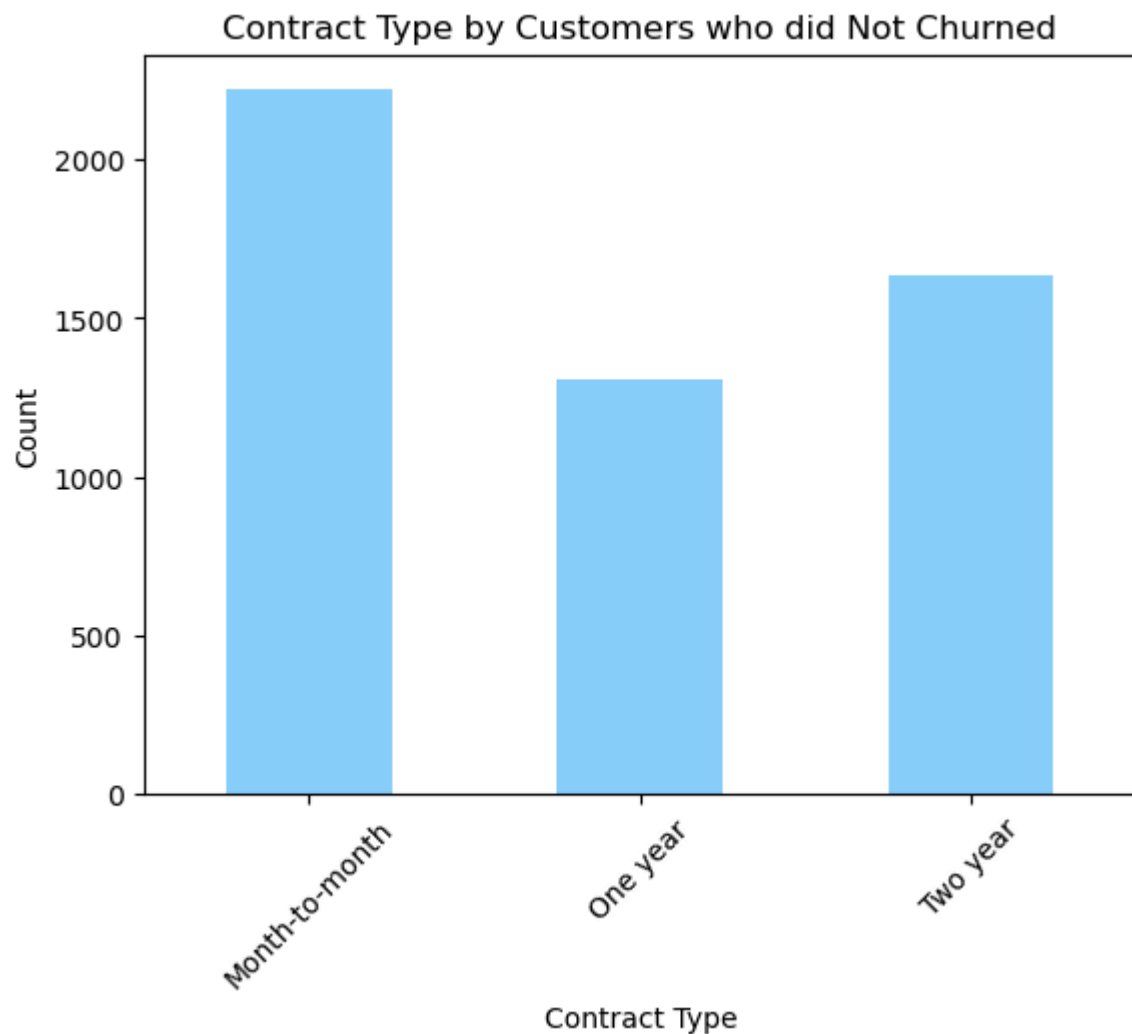
Contract Type by Customer who did not Churn

```
In [40]: #Customers who did not Churn
churn_counts = data[data['Churn'] == 'No'].groupby(['Contract']).size()
print(churn_counts)
```

```
Contract
Month-to-month    2220
One year          1306
Two year          1637
dtype: int64
```

```
In [41]: churn_counts = data[data['Churn'] == 'No'].groupby(['Contract']).size()

# Plotting
churn_counts.plot(kind='bar', color='lightskyblue')
plt.xlabel('Contract Type')
plt.ylabel('Count')
plt.title('Contract Type by Customers who did Not Churned')
plt.xticks(rotation=45)
plt.show()
```



Summary of the Contract Type

```
In [42]: import matplotlib.pyplot as plt
import matplotlib.ticker as mtick

colors = ['orange', 'lightskyblue']
contract_churn = data.groupby(['Contract', 'Churn']).size().unstack()


ax = (contract_churn.T * 100.0 / contract_churn.T.sum()).T.plot(kind='bar',
                                                                    width=0.3,
                                                                    stacked=True,
                                                                    rot=0,
                                                                    figsize=(10,6),
                                                                    color=colors)

ax.yaxis.set_major_formatter(mtick.PercentFormatter())
ax.legend(loc='best', prop={'size': 14}, title='Churn')
ax.set_ylabel('% Customers', size=14)
ax.set_title('Churn by Contract Type', size=14)

for p in ax.patches:
    width, height = p.get_width(), p.get_height()
    x, y = p.get_xy()
    ax.annotate('{:.0f}%'.format(height), (p.get_x() + 0.25 * width, p.get_y() + 0.4 *
                                                                    height),
                color='white',
                weight='bold',
```


```
size=14)
```

```
plt.show()
```




Distribution of Monthly Charges by Churn

```
In [43]: sns.set_context("paper", font_scale=1.1)
ax = sns.kdeplot(data.MonthlyCharges[(data["Churn"] == 'No') ],
                 color="orange", shade = True);
ax = sns.kdeplot(data.MonthlyCharges[(data["Churn"] == 'Yes') ],
                 ax=ax, color="Blue", shade= True);
ax.legend(["Not Churn", "Churn"], loc='upper right');
ax.set_ylabel('Density');
ax.set_xlabel('Monthly Charges');
ax.set_title('Distribution of Monthly Charges by Churn');
```



Distribution of Total Charges by Churn

```
In [44]: ax = sns.kdeplot(data.TotalCharges[(data["Churn"] == 'No') ],
                        color="Gold", shade = True);
ax = sns.kdeplot(data.TotalCharges[(data["Churn"] == 'Yes') ],
                ax=ax, color="Green", shade= True);
ax.legend(["Not Chuurn", "Churn"], loc='upper right');
ax.set_ylabel('Density');
ax.set_xlabel('Total Charges');
ax.set_title('Distribution of total charges by churn');
```

Tenure Vs Churn


```
In [45]: # Create the box plot
fig = px.box(data, x='Churn', y='tenure')

# Update y-axis properties
fig.update_yaxes(title_text='Tenure (Months)', row=1, col=1)
# Update x-axis properties
fig.update_xaxes(title_text='Churn', row=1, col=1)

# Update size and title
fig.update_layout(autosize=True, width=750, height=600,
                  title_font=dict(size=25, family='Courier'),
                  title='<b>Tenure vs Churn</b>')

fig.show()
```

Tenure vs Churn



The above shows that majority of new customers are likely to churn than old customer

Data Pre-processing

Here we;

- Further explored the relationships between variables and identified patterns or trends.
- Performed feature engineering by creating new features/columns and transforming existing ones to better represent the data and improve model performance.
- Utilised encoding techniques like Label Encoder and One Hot Encoder.
- Performed Normalization using Min Max Scaler
- Performed Feature Selection for dimensionality reduction using "Recursive Feature Elimination (RFE)" to reduce complexity and improve model performance.

Let's start with MultipleLines that has 3 observations and transform the "No Phone Service" to No as it also presents customers with no MultipleLines.

```
In [46]: data['MultipleLines'] = data['MultipleLines'].replace('No phone service', 'No')
data['MultipleLines'].value_counts()
```

```
Out[46]: No      4065
Yes       2967
Name: MultipleLines, dtype: int64
```

We will do the same for Online Security, OnlineBackup, DeviceProtection, TechSupport, StreamingTV, and StreamingMovies

```
In [47]: # Create a List for the columns to replace
columns_to_replace = ['OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies']

# Replace 'No internet service' in all specified columns with 'No'
data[columns_to_replace] = data[columns_to_replace].replace('No internet service', 'No')

# Check value counts for all specified columns after replacement
value_counts_all = data[columns_to_replace].apply(pd.value_counts)

# Print the value counts for all specified columns
print(value_counts_all)
```

	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV	\
No	5017	4607	4614	4992	4329	
Yes	2015	2425	2418	2040	2703	

	StreamingMovies
No	4301
Yes	2731

Label Encoding

Here we will perform Label Encoding to transform categorical columns with 2 observations to 0 and 1

And use One Hot Encoder for columns with more than 2 observations

```
In [48]: data.head()
```

Out[48]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService
0	Female	No	Yes	No	1	No	No	DSL
1	Male	No	No	No	34	Yes	No	DSL
2	Male	No	No	No	2	Yes	No	DSL
3	Male	No	No	No	45	No	No	DSL
4	Female	No	No	No	2	Yes	No	Fiber optic

We will exclude Columns 'Internet Service, Contract, Payment Method, Monthly Charges and Total Charges' we will treat them after as 'Internet Service, Contract, and Payment Method' have more than 2 outputs and Monthly Charges and Total Charges are already in numeric form

```
In [49]: # Get all column names
all_columns = data.columns

# Exclude the columns (Internet Service, Contract, Payment Method, Monthly Charges and
columns_to_exclude = ['InternetService', 'Contract', 'PaymentMethod', 'MonthlyCharges']

# Get the columns to encode by removing the excluded columns from all columns
columns_to_encode = [col for col in all_columns if col not in columns_to_exclude]

# Initialize LabelEncoder
label_encoder = LabelEncoder()

# Encode each column
data[columns_to_encode] = data[columns_to_encode].apply(label_encoder.fit_transform)
```

```
In [50]: data.dtypes
```

```
Out[50]: gender                int32
SeniorCitizen                int32
Partner                      int32
Dependents                   int32
tenure                       int64
PhoneService                 int32
MultipleLines                int32
InternetService              object
OnlineSecurity               int32
OnlineBackup                 int32
DeviceProtection             int32
TechSupport                  int32
StreamingTV                  int32
StreamingMovies              int32
Contract                     object
PaperlessBilling             int32
PaymentMethod                object
MonthlyCharges               int64
TotalCharges                 int64
Churn                        int32
dtype: object
```

Using One Hot Encoder

Now let's convert Internet Service, Contract, and Payment Method to numeric using One Hot Encoder

```
In [51]: # Perform one-hot encoding
data = pd.get_dummies(data, columns=['InternetService', 'Contract', 'PaymentMethod'])
```

```
In [52]: #Let's view the data to ensure the columns have been converted accordingly
data.head()
```

```
Out[52]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
1	1	0	0	0	33	1	0	1	0	0	0	0	0	0	0	0	0	0	0
2	1	0	0	0	1	1	0	1	0	0	0	0	0	0	0	0	0	0	0
3	1	0	0	0	44	0	0	0	0	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0

5 rows × 27 columns

```
In [53]: data.dtypes
```

```
Out[53]: gender                int32
SeniorCitizen                int32
Partner                      int32
Dependents                   int32
tenure                       int64
PhoneService                 int32
MultipleLines                int32
OnlineSecurity               int32
OnlineBackup                 int32
DeviceProtection             int32
TechSupport                  int32
StreamingTV                  int32
StreamingMovies              int32
PaperlessBilling             int32
MonthlyCharges               int64
TotalCharges                  int64
Churn                        int32
InternetService_DSL          uint8
InternetService_Fiber optic  uint8
InternetService_No           uint8
Contract_Month-to-month      uint8
Contract_One year            uint8
Contract_Two year            uint8
PaymentMethod_Bank transfer (automatic) uint8
PaymentMethod_Credit card (automatic)  uint8
PaymentMethod_Electronic check         uint8
PaymentMethod_Mailed check             uint8
dtype: object
```

```
In [54]: data.head()
```

```
Out[54]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	OnlineSecurity	Churn
0	0	0	1	0	0	0	0	0	0
1	1	0	0	0	33	1	0	1	1
2	1	0	0	0	1	1	0	1	1
3	1	0	0	0	44	0	0	1	1
4	0	0	0	0	1	1	0	0	0

5 rows × 27 columns

```
In [55]: data.shape
```

```
Out[55]: (7032, 27)
```

Correlation Analysis

```
In [56]: # Calculate correlations between all columns and the target variable "Churn"
correlation_with_churn = data.corr()['Churn'].sort_values(ascending=False)

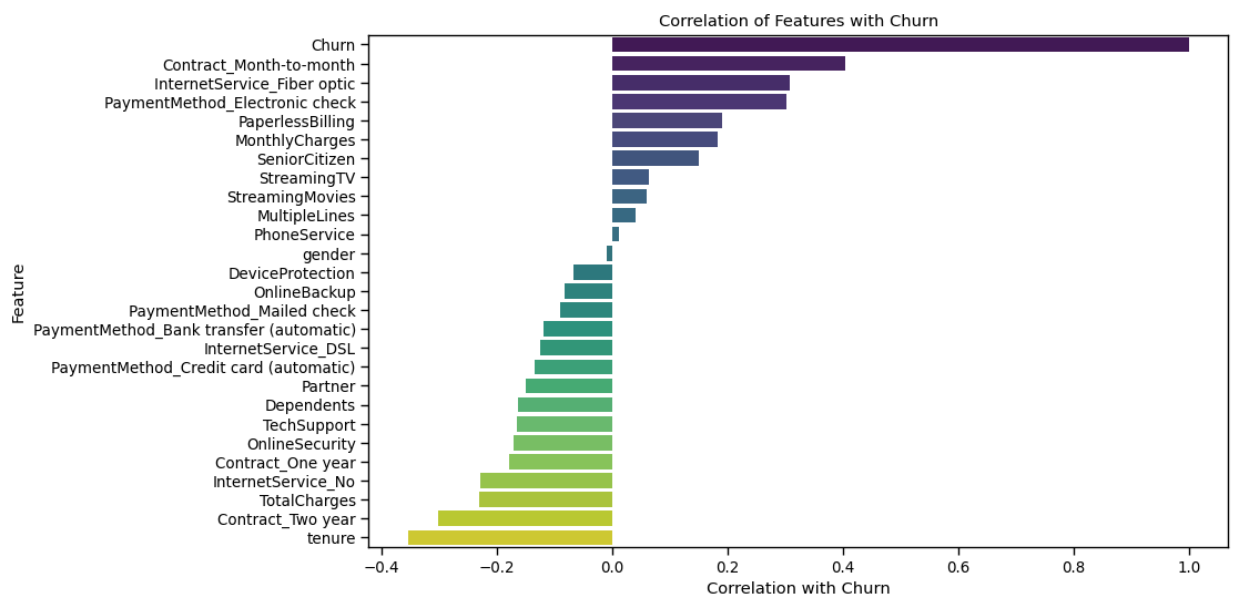
# Print correlation values
print(correlation_with_churn)
```

Churn	1.000000
Contract_Month-to-month	0.404565
InternetService_Fiber optic	0.307463
PaymentMethod_Electronic check	0.301455
PaperlessBilling	0.191454
MonthlyCharges	0.182989
SeniorCitizen	0.150541
StreamingTV	0.063254
StreamingMovies	0.060860
MultipleLines	0.040033
PhoneService	0.011691
gender	-0.008545
DeviceProtection	-0.066193
OnlineBackup	-0.082307
PaymentMethod_Mailed check	-0.090773
PaymentMethod_Bank transfer (automatic)	-0.118136
InternetService_DSL	-0.124141
PaymentMethod_Credit card (automatic)	-0.134687
Partner	-0.149982
Dependents	-0.163128
TechSupport	-0.164716
OnlineSecurity	-0.171270
Contract_One year	-0.178225
InternetService_No	-0.227578
TotalCharges	-0.230843
Contract_Two year	-0.301552
tenure	-0.354049

Name: Churn, dtype: float64

```
In [57]: # Calculate correlations between all columns and the target variable "Churn"
correlation_with_churn = data.corr()['Churn'].sort_values(ascending=False)

# Plot the correlation values
plt.figure(figsize=(10, 6))
sns.barplot(x=correlation_with_churn.values, y=correlation_with_churn.index, palette='
plt.xlabel('Correlation with Churn')
plt.ylabel('Feature')
plt.title('Correlation of Features with Churn')
plt.show()
```



From the above we can see that there is no perfect correlation with the target variable

Feature Selection Using Recursive Feature Elimination Technique

Since we have 26 columns with no perfect correlation to the target variable we will be using a feature selection techniques to identify the most relevant variables for building the model. This would helps to improve the model performance, reduce overfitting, and enhance interpretability. For this we will be using 'Recursive Feature Elimination' (RFE) to help identify important features. It works by recursively removing features and builds the model on the remaining features until a specified number of features is reached.

Model Building and Evaluation.

- Choose appropriate statistical or machine learning models based on the problem and data characteristics.
- Split the data into training and testing sets for model evaluation.
- Train the models on the training data and evaluate their performance using appropriate metrics.
- Fine-tune model parameters and compare different models to select the best-performing one.
- Iterate and Refine Review the analysis process and results, and iterate as needed to refine the analysis or address new questions or insights.

Splitting the Data into Training and Testing

```
In [58]: X = data.drop('Churn', axis = 1)

y = data['Churn']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=
```

Normalizing the Data Using MinMaxScaler

We will apply normalization technique on the dataset because the dataset columns are of different scale. Here we will using the Min Max Scaler to normalize the dataset.

```
In [59]: # Initialize the MinMaxScaler
scaler = MinMaxScaler()
```

```
In [60]: # Fit the scaler to your training data and transform it
X_train_normalized = scaler.fit_transform(X_train)
```



```
# Transform the test data using the same scaler
X_test_normalized = scaler.transform(X_test)

# Now X_train_normalized and X_test_normalized contain the normalized data
```

Selecting the Model

Here we will be using 4 models

- Logistic Regression Model
- Random Forest Classifier
- Support Vector Machine Classifier

After which we will select the best fit model

Models

```
In [61]: #Initializing Models
logistic_regression = LogisticRegression()
random_forest = RandomForestClassifier()
svm_model = SVC(probability=True)
```

EXPERIMENT 1

Using All the Features without RFE

First let's build our model using all the features after which we will use the RFE on top 10 and 20

Logistic Regression Model

```
In [62]: # Train and evaluate Logistic Regression model on normalized data
logistic_regression.fit(X_train_normalized, y_train)
y_pred_lr_normalized = logistic_regression.predict(X_test_normalized)
accuracy_lr_normalized = accuracy_score(y_test, y_pred_lr_normalized)

print("Logistic Regression Accuracy (Normalized Data - Test Set):", accuracy_lr_normalized)

# Predict on the training set
y_pred_lr_normalized_train = logistic_regression.predict(X_train_normalized)

# Compute accuracy for the training set
accuracy_lr_normalized_train = accuracy_score(y_train, y_pred_lr_normalized_train)

# Print accuracy for the training set
print("Logistic Regression Accuracy (Normalized Data - Training Set):", accuracy_lr_normalized_train)

print("Logistic Regression Classification Report (Normalized Data):")
print(classification_report(y_test, y_pred_lr_normalized))
```

```
conf_matrix_lr_normalized = confusion_matrix(y_test, y_pred_lr_normalized)
print("Logistic Regression Confusion Matrix (Normalized Data):")
print(conf_matrix_lr_normalized)
```


Logistic Regression Accuracy (Normalized Data - Test Set): 0.7924662402274343
 Logistic Regression Accuracy (Normalized Data - Training Set): 0.8117333333333333
 Logistic Regression Classification Report (Normalized Data):

	precision	recall	f1-score	support
0	0.83	0.91	0.87	1033
1	0.65	0.48	0.55	374
accuracy			0.79	1407
macro avg	0.74	0.69	0.71	1407
weighted avg	0.78	0.79	0.78	1407

Logistic Regression Confusion Matrix (Normalized Data):
 [[936 97]
 [195 179]]

```
In [63]: # Plotting the Confusion Matrix for Logistic Regression Model with Normalized Data
cm_display_lr_normalized = ConfusionMatrixDisplay(confusion_matrix=conf_matrix_lr_norm
cm_display_lr_normalized.plot()
plt.title("Logistic Regression Confusion Matrix Using All Features Without RFE (Normal
plt.show()
```


Logistic Regression Confusion Matrix Using All Features Without RFE (Normalized Data)



```
In [64]: # Plotting the ROC curve for Logistic Regression Model with Normalized Data
y_prob_lr_normalized = logistic_regression.predict_proba(X_test_normalized)[: , 1]
fpr_lr_normalized, tpr_lr_normalized, thresholds_lr_normalized = roc_curve(y_test, y_p
auc_lr_normalized = roc_auc_score(y_test, y_prob_lr_normalized)

plt.figure(figsize=(8, 6))
plt.plot(fpr_lr_normalized, tpr_lr_normalized, label='ROC curve (area = %0.2f)' % auc_
```

```
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```



SVM Model

```
In [65]: # Train the SVM model on normalized data without RFE
svm_model.fit(X_train_normalized, y_train)

# Make predictions
y_pred = svm_model.predict(X_test_normalized)

# Make predictions
y_pred_train = svm_model.predict(X_train_normalized)

# Compute the accuracy of the train set
train_accuracy = accuracy_score(y_train, y_pred_train)
print("Accuracy on the train set:", train_accuracy)

# Evaluate the model
print("Scenario: Normalized data without RFE")
```

```

print("Classification Report:")
print(classification_report(y_test, y_pred))

# Compute the confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(conf_matrix)

```

Accuracy on the train set: 0.8241777777777778

Scenario: Normalized data without RFE

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.89	0.86	1033
1	0.62	0.48	0.54	374
accuracy			0.78	1407
macro avg	0.72	0.69	0.70	1407
weighted avg	0.77	0.78	0.78	1407

Confusion Matrix:

```

[[923 110]
 [193 181]]

```


In [66]:

```

# Plot the confusion matrix
cm_display = ConfusionMatrixDisplay(confusion_matrix=conf_matrix)
cm_display.plot()
plt.title("Confusion Matrix - Normalized data Using All Features Without RFE")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()

```

Confusion Matrix - Normalized data Using All Features Without RFE



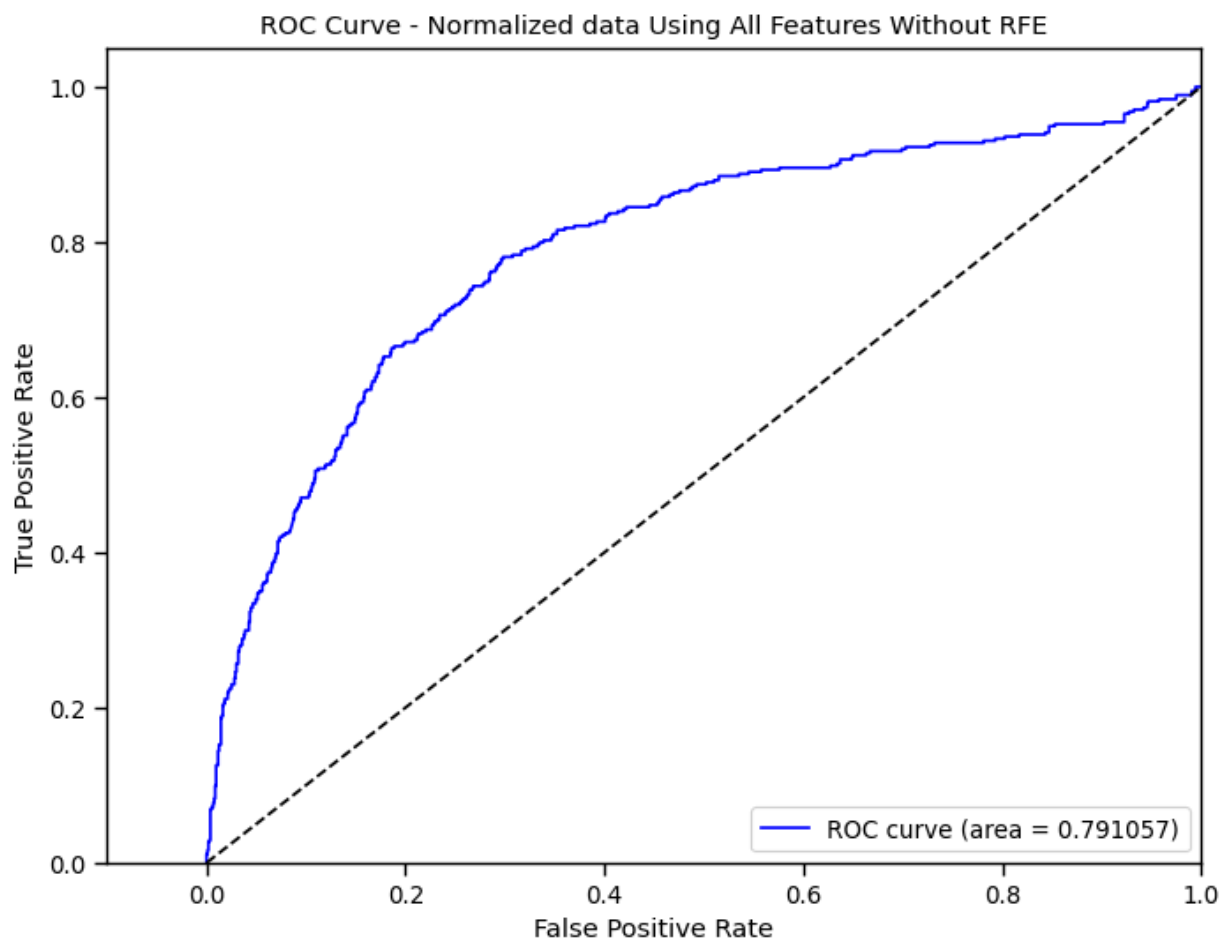
```
In [67]: # Compute the probability estimates for positive class
y_prob = svm_model.predict_proba(X_test_normalized)[: , 1]

# Compute fpr, tpr, and thresholds
fpr, tpr, thresholds = roc_curve(y_test, y_prob)

# Compute AUC
auc = roc_auc_score(y_test, y_prob)
print('AUC:', auc)

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', label='ROC curve (area = %f)' % auc)
plt.plot([0, 1], [0, 1], linestyle='--', color='black')
plt.xlim([-0.1, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc="lower right")
plt.title('ROC Curve - Normalized data Using All Features Without RFE')
plt.show()
```

AUC: 0.7910568874209898



Random Forest Model

```
In [68]: # Train and evaluate Random Forest model on normalized data
random_forest.fit(X_train_normalized, y_train)
y_pred_rf = random_forest.predict(X_test_normalized)

# Make predictions on the train set
y_pred_train_rf = random_forest.predict(X_train_normalized)

# Compute the accuracy of the train set for Random Forest model
train_accuracy_rf = accuracy_score(y_train, y_pred_train_rf)
print("Random Forest Accuracy on the train set (Normalized Data):", train_accuracy_rf)

# Compute confusion matrix for Random Forest model
conf_matrix_rf = confusion_matrix(y_test, y_pred_rf)

# Print accuracy, classification report, and confusion matrix
accuracy_rf = accuracy_score(y_test, y_pred_rf)
print("Random Forest Accuracy (Normalized Data - Test Set):", accuracy_rf)
print("Random Forest Classification Report (Normalized Data):")
print(classification_report(y_test, y_pred_rf))
print("Random Forest Confusion Matrix (Normalized Data):")
print(conf_matrix_rf)
```

Random Forest Accuracy on the train set (Normalized Data): 0.9976888888888888

Random Forest Accuracy (Normalized Data - Test Set): 0.7789623312011372

Random Forest Classification Report (Normalized Data):

	precision	recall	f1-score	support
0	0.83	0.88	0.85	1033
1	0.60	0.50	0.54	374
accuracy			0.78	1407
macro avg	0.72	0.69	0.70	1407
weighted avg	0.77	0.78	0.77	1407


Random Forest Confusion Matrix (Normalized Data):

```
[[910 123]
```

```
[188 186]]
```

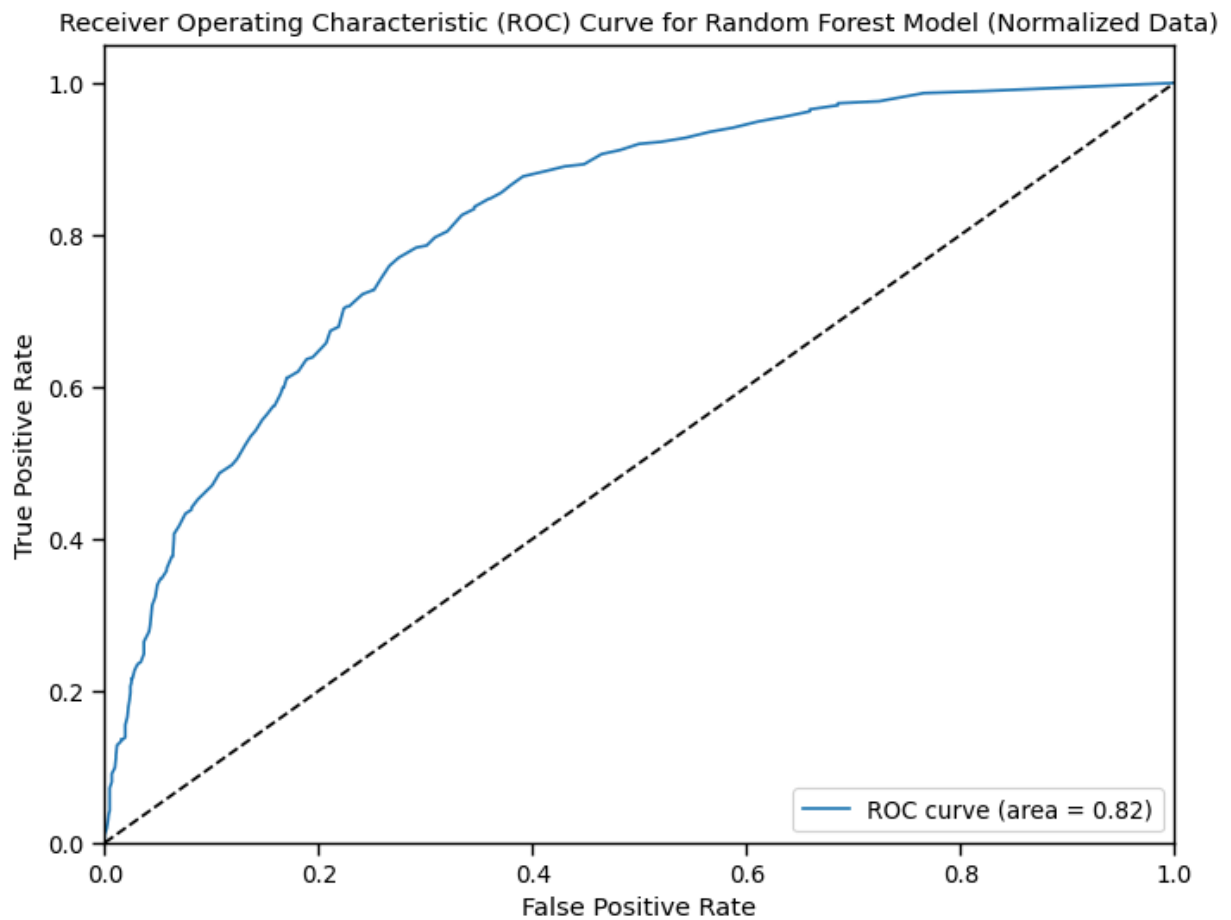
```
In [69]: # Plotting the Confusion Matrix for Random Forest Model with Normalized Data
cm_display_rf = ConfusionMatrixDisplay(confusion_matrix=conf_matrix_rf, display_labels=
cm_display_rf.plot()
plt.title("Random Forest Confusion Matrix Using All Features Without RFE (Normalized D
plt.show())
```

Random Forest Confusion Matrix Using All Features Without RFE (Normalized Data)



```
In [70]: # Plotting the ROC curve for Random Forest Model with Normalized Data
y_prob_rf = random_forest.predict_proba(X_test_normalized)[: , 1]
fpr_rf, tpr_rf, thresholds_rf = roc_curve(y_test, y_prob_rf)
auc_rf = roc_auc_score(y_test, y_prob_rf)

plt.figure(figsize=(8, 6))
plt.plot(fpr_rf, tpr_rf, label='ROC curve (area = %0.2f)' % auc_rf)
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve for Random Forest Model (Norm
plt.legend(loc="lower right")
plt.show()
```



EXPERIMENT 2

Using RFE on Logistic Regression, Random Forest and SVM to Select the Top 10 Features.

```
In [71]: # Function to perform Recursive Feature Elimination (RFE)
def perform_rfe(model, X_train_normalized, y_train, n_features):
    rfe = RFE(model, n_features_to_select=n_features)
    rfe.fit(X_train_normalized, y_train)
    selected_features = rfe.support_
    return selected_features
```

```
In [72]: # Initialize models
logistic_regression = LogisticRegression(max_iter=1000)
svm_model_linear = SVC(kernel='linear', probability=True)
random_forest = RandomForestClassifier()
```

Logistic Regression Model

```
In [73]: # Get Selected Features for each model using RFE with normalized data
logistic_regression_features = perform_rfe(logistic_regression, X_train_normalized, y_
```



```
In [74]: # Train and evaluate Logistic Regression model with normalized data
logistic_regression.fit(X_train_normalized[:, logistic_regression_features], y_train)
y_pred_lr = logistic_regression.predict(X_test_normalized[:, logistic_regression_features])
accuracy_lr = accuracy_score(y_test, y_pred_lr)

# Make predictions on the train set
y_pred_train_lr = logistic_regression.predict(X_train_normalized[:, logistic_regression_features])

# Compute the accuracy of the train set for Logistic Regression model
train_accuracy_lr = accuracy_score(y_train, y_pred_train_lr)
print("Logistic Regression Accuracy on the train set:", train_accuracy_lr)

# compute the confusion matrix
conf_matrix_lr = confusion_matrix(y_test, y_pred_lr)

# Print accuracy and classification report
print("Logistic Regression Accuracy - Test Set:", accuracy_lr)
print("Logistic Regression Classification Report:")
print(classification_report(y_test, y_pred_lr))
```

Logistic Regression Accuracy on the train set: 0.8030222222222222


Logistic Regression Accuracy - Test Set: 0.7903340440653873

Logistic Regression Classification Report:

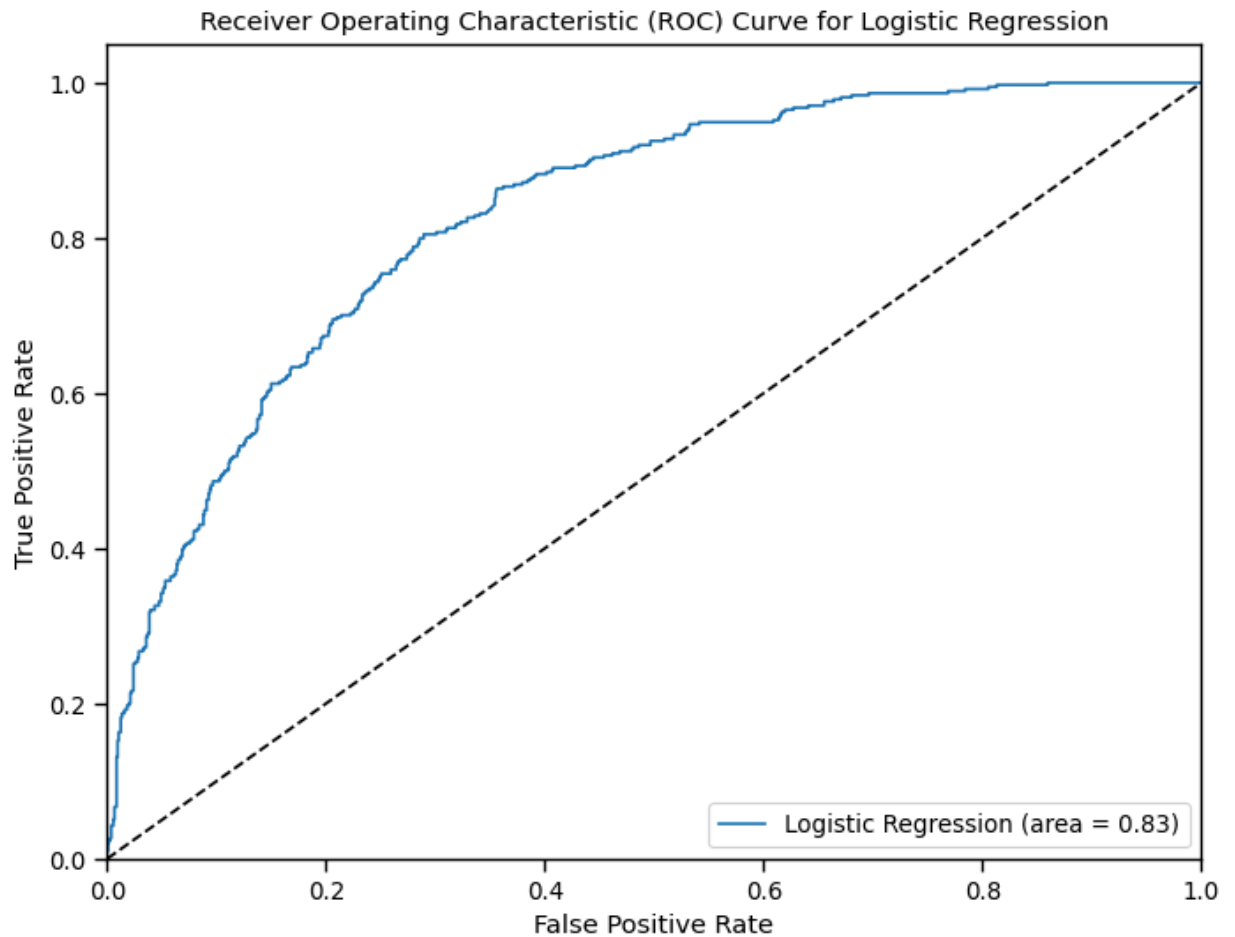
	precision	recall	f1-score	support
0	0.83	0.91	0.86	1033
1	0.64	0.47	0.54	374
accuracy			0.79	1407
macro avg	0.74	0.69	0.70	1407
weighted avg	0.78	0.79	0.78	1407

```
In [75]: # Plot the Confusion Matrix for Logistic Regression Model with Normalized Data
cm_display_lr = ConfusionMatrixDisplay(confusion_matrix=conf_matrix_lr)
cm_display_lr.plot()
plt.title("Logistic Regression Confusion Matrix Using RFE Top 10 Features (Normalized Data)")
plt.show()
```

Logistic Regression Confusion Matrix Using RFE Top 10 Features (Normalized Data)



```
In [76]: # Plot ROC AUC curve for Logistic Regression
y_prob_lr = logistic_regression.predict_proba(X_test_normalized[:, logistic_regression
fpr_lr, tpr_lr, thresholds_lr = roc_curve(y_test, y_prob_lr)
auc_lr = roc_auc_score(y_test, y_prob_lr)
plt.figure(figsize=(8, 6))
plt.plot(fpr_lr, tpr_lr, label='Logistic Regression (area = %0.2f)' % auc_lr)
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve for Logistic Regression')
plt.legend(loc="lower right")
plt.show()
```



```
In [77]: # This function to prints the columns selected by RFE
#def print_selected_columns(selected_features):
#    # Print the indices of the selected features

#    #print("Indices of selected columns:")
#    #for index, selected in enumerate(selected_features):
#    #    # if selected:
#    #        print(index)

# Assuming you have obtained the selected features from RFE
# For example, logistic_regression_features from your code snippet

# Print indices of selected columns
#print_selected_columns(logistic_regression_features)
```

Random Forest Model

```
In [78]: # Get Selected Features for each model using RFE with normalized data
random_forest_features = perform_rfe(random_forest, X_train_normalized, y_train, n_features_to_select=10)
```

```
In [79]: # Train and evaluate Random Forest model with normalized data
random_forest.fit(X_train_normalized[:, random_forest_features], y_train)
y_pred_rf = random_forest.predict(X_test_normalized[:, random_forest_features])
```

```

accuracy_rf = accuracy_score(y_test, y_pred_rf)

# Make predictions on the train set
y_pred_train_rf = random_forest.predict(X_train_normalized[:, random_forest_features])

# Compute the accuracy of the train set for Random Forest model
train_accuracy_rf = accuracy_score(y_train, y_pred_train_rf)
print("Random Forest Accuracy on the train set:", train_accuracy_rf)

# Compute the confusion matrix
conf_matrix_rf = confusion_matrix(y_test, y_pred_rf)

# Print accuracy and classification report
print("Random Forest Accuracy:", accuracy_rf)
print("Random Forest Classification Report:")
print(classification_report(y_test, y_pred_rf))

```

Random Forest Accuracy on the train set: 0.9971555555555556

Random Forest Accuracy: 0.775408670931059

Random Forest Classification Report:

	precision	recall	f1-score	support
0	0.82	0.89	0.85	1033
1	0.60	0.46	0.52	374
accuracy			0.78	1407
macro avg	0.71	0.67	0.69	1407
weighted avg	0.76	0.78	0.77	1407


In [80]:

```

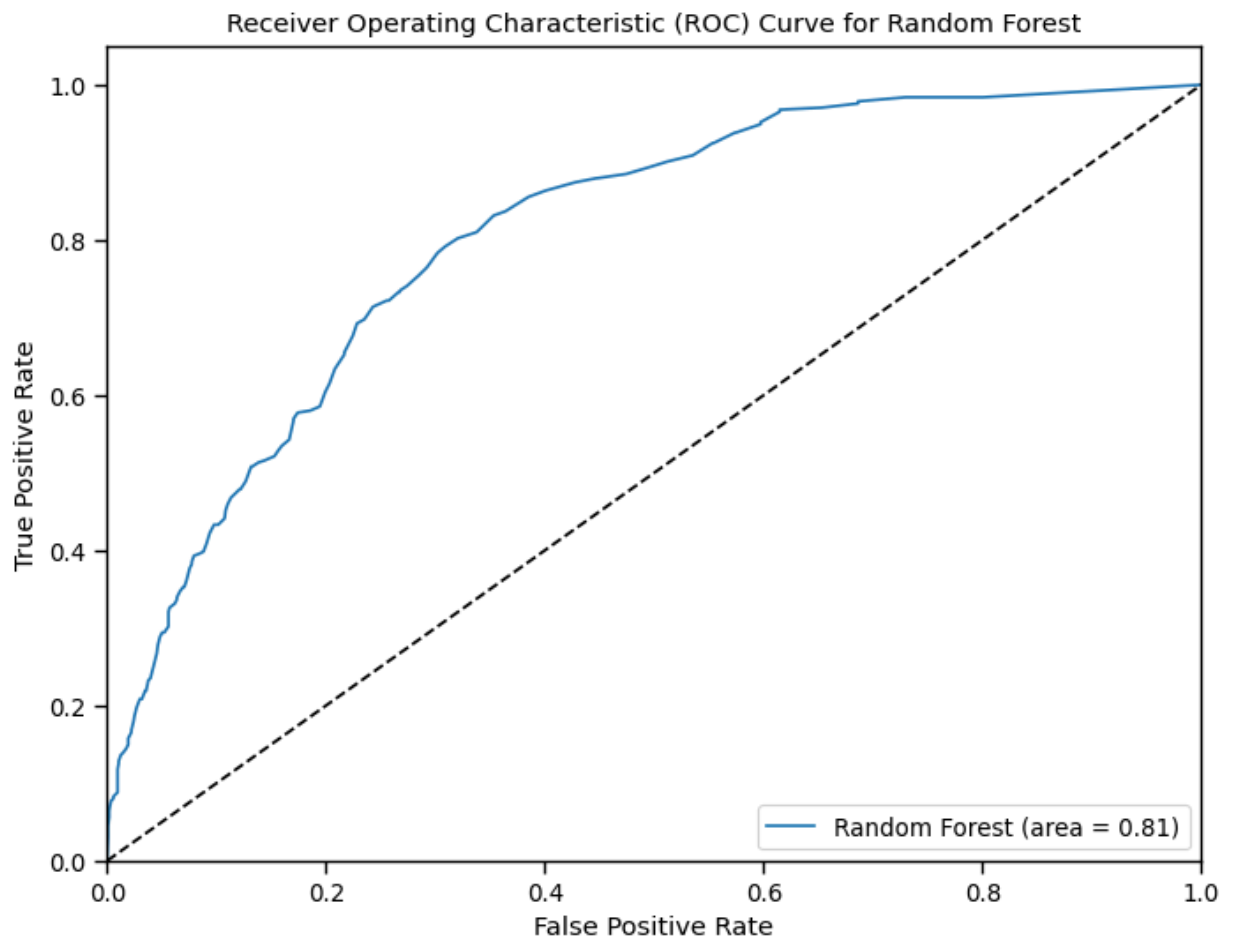
# Plot the Confusion Matrix for Random Forest Model with Normalized Data
cm_display_rf = ConfusionMatrixDisplay(confusion_matrix=conf_matrix_rf)
cm_display_rf.plot()
plt.title("Random Forest Confusion Matrix Using RFE Top 10 Features (Normalized Data)")
plt.show()

```

Random Forest Confusion Matrix Using RFE Top 10 Features (Normalized Data)



```
In [81]: # Plot ROC AUC curve for Random Forest
y_prob_rf = random_forest.predict_proba(X_test_normalized[:, random_forest_features])[
fpr_rf, tpr_rf, thresholds_rf = roc_curve(y_test, y_prob_rf)
auc_rf = roc_auc_score(y_test, y_prob_rf)
plt.figure(figsize=(8, 6))
plt.plot(fpr_rf, tpr_rf, label='Random Forest (area = %0.2f)' % auc_rf)
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve for Random Forest')
plt.legend(loc="lower right")
plt.show()
```



SVM Model

```
In [82]: # Initialize RFE for feature selection with SVM model
rfe = RFE(estimator=svm_model_linear, n_features_to_select=10)

# Fit RFE on normalized data
rfe.fit(X_train_normalized, y_train)

# Get the selected features
selected_features_svm = rfe.support_

# Select top 10 features from training and test data
X_train_top10_svm = X_train_normalized[:, selected_features_svm]
X_test_top10_svm = X_test_normalized[:, selected_features_svm]
```

```
In [83]: # Train the SVM model on top 10 features
svm_model_linear.fit(X_train_top10_svm, y_train)

# Make predictions on the test set
y_pred_svm = svm_model_linear.predict(X_test_top10_svm)

# Make predictions on the train set
y_pred_train_svm = svm_model_linear.predict(X_train_top10_svm)

# Compute the accuracy of the train set for SVM model
train_accuracy_svm = accuracy_score(y_train, y_pred_train_svm)
```

```
print("SVM Accuracy on the train set with top 10 features:", train_accuracy_svm)

# Evaluate the model
print("Scenario: Normalized data with RFE on top 10 features (SVM with linear kernel)")
print("Classification Report:")
print(classification_report(y_test, y_pred_svm))
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred_svm))
```

SVM Accuracy on the train set with top 10 features: 0.8

Scenario: Normalized data with RFE on top 10 features (SVM with linear kernel)

Classification Report:

	precision	recall	f1-score	support
0	0.81	0.94	0.87	1033
1	0.69	0.38	0.49	374
accuracy			0.79	1407
macro avg	0.75	0.66	0.68	1407
weighted avg	0.78	0.79	0.77	1407


Confusion Matrix:

```
[[969  64]
```

```
 [232 142]]
```

```
In [84]: # Plot the Confusion Matrix for SVM Model with Linear Kernel and RFE-selected features
cm_display_svm = ConfusionMatrixDisplay(confusion_matrix=confusion_matrix(y_test, y_pr
cm_display_svm.plot()
plt.title("Confusion Matrix - Normalized data with RFE on top 10 features (SVM with li
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()
```

Confusion Matrix - Normalized data with RFE on top 10 features (SVM with linear kernel)



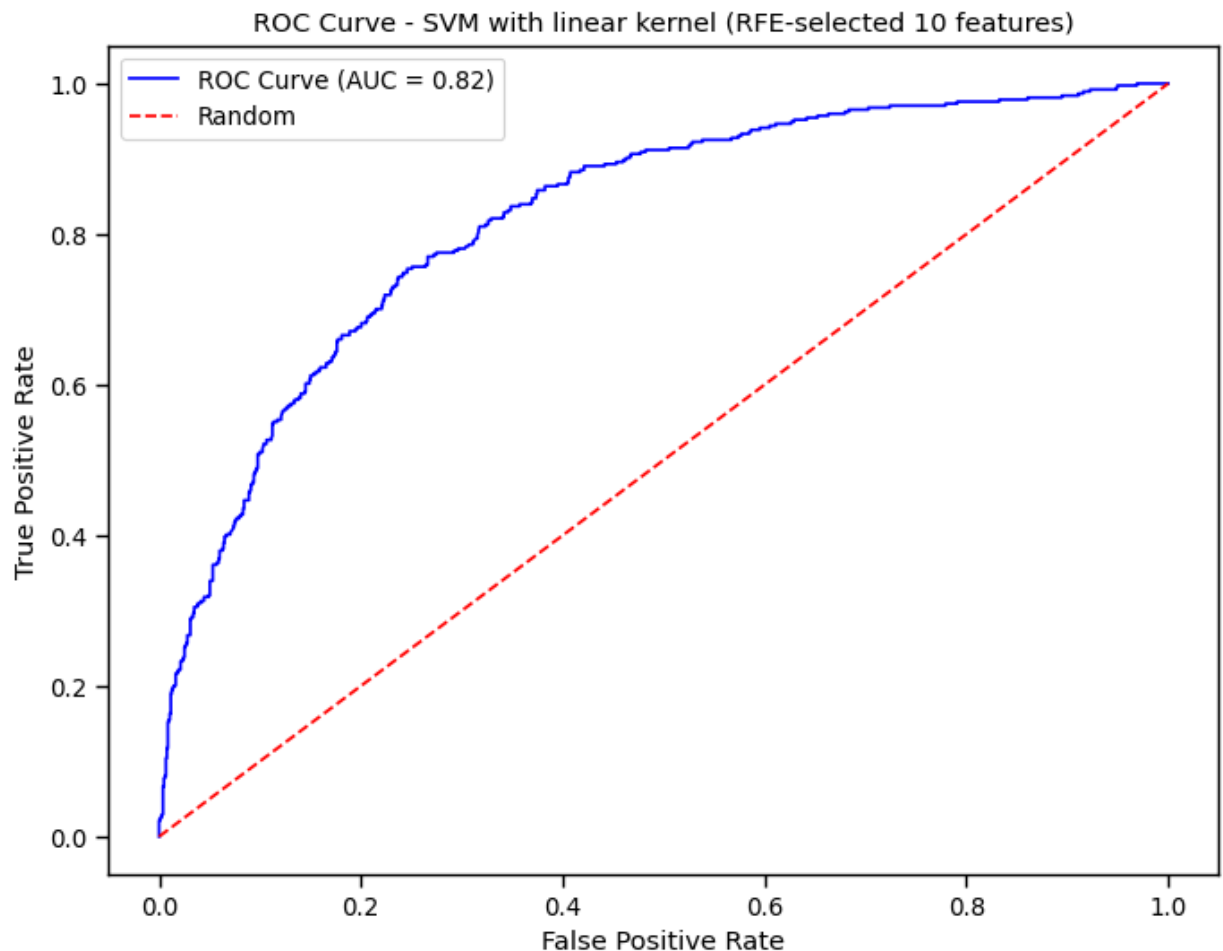
```
In [85]: # Compute the probability estimates for positive class
y_prob_svm = svm_model_linear.predict_proba(X_test_top10_svm)[: , 1]

# Compute fpr, tpr, and thresholds
fpr_svm, tpr_svm, thresholds_svm = roc_curve(y_test, y_prob_svm)

# Compute ROC AUC score
roc_auc_svm = roc_auc_score(y_test, y_prob_svm)
print('ROC AUC (SVM with linear kernel):', roc_auc_svm)

# Plot ROC AUC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr_svm, tpr_svm, color='blue', label='ROC Curve (AUC = {:.2f})'.format(roc_auc_svm))
plt.plot([0, 1], [0, 1], color='red', linestyle='--', label='Random')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - SVM with linear kernel (RFE-selected 10 features)')
plt.legend()
plt.show()
```

ROC AUC (SVM with linear kernel): 0.8217097804535878



EXPERIMENT 3

PERFORMING RFE ON TOP TWENTY (20) FEATURES

Logistic Regression, SVM Model and Random Forest Using 20 Top features

Models

Logistics Regression Model

```
In [86]: # Get Selected Features for each model using RFE with normalized data
logistic_regression_features = perform_rfe(logistic_regression, X_train_normalized, y_
```

```
In [87]: # Train and evaluate Logistic Regression model with normalized data
logistic_regression.fit(X_train_normalized[:, logistic_regression_features], y_train)
y_pred_lr = logistic_regression.predict(X_test_normalized[:, logistic_regression_featu
accuracy_lr = accuracy_score(y_test, y_pred_lr)

# Make predictions on the train set
y_pred_train_lr = logistic_regression.predict(X_train_normalized[:, logistic_regressio

# Compute the accuracy of the train set for Logistic Regression model
train_accuracy_lr = accuracy_score(y_train, y_pred_train_lr)
print("Logistic Regression Accuracy on the train set:", train_accuracy_lr)

# Compute the confusion matrix
conf_matrix_lr = confusion_matrix(y_test, y_pred_lr)

# Print accuracy and classification report
print("Logistic Regression Accuracy:", accuracy_lr)
print("Logistic Regression Classification Report:")
print(classification_report(y_test, y_pred_lr))
```

Logistic Regression Accuracy on the train set: 0.8094222222222223


Logistic Regression Accuracy: 0.7889125799573561

Logistic Regression Classification Report:


	precision	recall	f1-score	support
0	0.83	0.90	0.86	1033
1	0.64	0.48	0.55	374
accuracy			0.79	1407
macro avg	0.73	0.69	0.70	1407
weighted avg	0.78	0.79	0.78	1407

```
In [88]: # Plot the Confusion Matrix for Logistic Regression Model with Normalized Data
cm_display_lr = ConfusionMatrixDisplay(confusion_matrix=conf_matrix_lr)
cm_display_lr.plot()
plt.title("Logistic Regression Confusion Matrix Using RFE Top 10 Features (Normalized
plt.show()
```

Logistic Regression Confusion Matrix Using RFE Top 10 Features (Normalized Data)



```
In [89]: # Plot ROC AUC curve for Logistic Regression
y_prob_lr = logistic_regression.predict_proba(X_test_normalized[:, logistic_regression
fpr_lr, tpr_lr, thresholds_lr = roc_curve(y_test, y_prob_lr)
auc_lr = roc_auc_score(y_test, y_prob_lr)
plt.figure(figsize=(8, 6))
plt.plot(fpr_lr, tpr_lr, label='Logistic Regression (area = %0.2f)' % auc_lr)
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve for Logistic Regression')
plt.legend(loc="lower right")
plt.show()
```



Random Forest Model

```
In [90]: # Get Selected Features for each model using RFE with normalized data
random_forest_features = perform_rfe(random_forest, X_train_normalized, y_train, n_features_to_select=10)
```

```
In [91]: # Train and evaluate Random Forest model with normalized data
random_forest.fit(X_train_normalized[:, random_forest_features], y_train)
y_pred_rf = random_forest.predict(X_test_normalized[:, random_forest_features])
accuracy_rf = accuracy_score(y_test, y_pred_rf)

# Make predictions on the train set
y_pred_train_rf = random_forest.predict(X_train_normalized[:, random_forest_features])

# Compute the accuracy of the train set for Random Forest model
train_accuracy_rf = accuracy_score(y_train, y_pred_train_rf)
print("Random Forest Accuracy on the train set:", train_accuracy_rf)

# Compute the confusion matrix
conf_matrix_rf = confusion_matrix(y_test, y_pred_rf)

# Print accuracy and classification report
print("Random Forest Accuracy:", accuracy_rf)
```

```
print("Random Forest Classification Report:")
print(classification_report(y_test, y_pred_rf))
```

Random Forest Accuracy on the train set: 0.9976888888888888


Random Forest Accuracy: 0.7917555081734187

Random Forest Classification Report:

	precision	recall	f1-score	support
0	0.83	0.90	0.86	1033
1	0.65	0.48	0.55	374
accuracy			0.79	1407
macro avg	0.74	0.69	0.71	1407
weighted avg	0.78	0.79	0.78	1407


```
In [92]: # Plot the Confusion Matrix for Random Forest Model with Normalized Data
cm_display_rf = ConfusionMatrixDisplay(confusion_matrix=conf_matrix_rf)
cm_display_rf.plot()
plt.title("Random Forest Confusion Matrix Using RFE Top 10 Features (Normalized Data)")
plt.show()
```

Random Forest Confusion Matrix Using RFE Top 10 Features (Normalized Data)



```
In [93]: # Plot ROC AUC curve for Random Forest
y_prob_rf = random_forest.predict_proba(X_test_normalized[:, random_forest_features])
fpr_rf, tpr_rf, thresholds_rf = roc_curve(y_test, y_prob_rf)
auc_rf = roc_auc_score(y_test, y_prob_rf)
plt.figure(figsize=(8, 6))
plt.plot(fpr_rf, tpr_rf, label='Random Forest (area = %0.2f)' % auc_rf)
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
```

```
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve for Random Forest')
plt.legend(loc="lower right")
plt.show()
```



SVM Model

```
In [94]: # Initialize RFE for feature selection with SVM model
rfe = RFE(estimator=svm_model_linear, n_features_to_select=20)
```

```
In [95]: # Fit RFE on normalized data
rfe.fit(X_train_normalized, y_train)

# Get the selected features
selected_features_svm = rfe.support_

# Select top 20 features from training and test data
X_train_top20_svm = X_train_normalized[:, selected_features_svm]
X_test_top20_svm = X_test_normalized[:, selected_features_svm]
```

```
In [96]: # Train the SVM model on top 20 features
svm_model_linear.fit(X_train_top20_svm, y_train)

# Make predictions on the test set
```

```

y_pred_svm = svm_model_linear.predict(X_test_top20_svm)

# Make predictions on the train set
y_pred_train_svm = svm_model_linear.predict(X_train_top20_svm)

# Compute the accuracy of the train set for SVM model
train_accuracy_svm = accuracy_score(y_train, y_pred_train_svm)
print("SVM Accuracy on the train set with top 20 features:", train_accuracy_svm)

# Evaluate the model
print("Scenario: Normalized data with RFE on top 20 features (SVM with linear kernel)")
print("Classification Report:")
print(classification_report(y_test, y_pred_svm))
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred_svm))

```

SVM Accuracy on the train set with top 20 features: 0.8062222222222222
Scenario: Normalized data with RFE on top 20 features (SVM with linear kernel)
Classification Report:

	precision	recall	f1-score	support
0	0.82	0.93	0.87	1033
1	0.69	0.42	0.52	374
accuracy			0.79	1407
macro avg	0.75	0.68	0.70	1407
weighted avg	0.78	0.79	0.78	1407

Confusion Matrix:
[[961 72]
[217 157]]


In [97]:

```

# Plot the Confusion Matrix for SVM Model with Linear Kernel and RFE-selected features
cm_display_svm = ConfusionMatrixDisplay(confusion_matrix=confusion_matrix(y_test, y_pr
cm_display_svm.plot()
plt.title("Confusion Matrix - Normalized data with RFE on top 20 features (SVM with li
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()

```

Confusion Matrix - Normalized data with RFE on top 20 features (SVM with linear kernel)




```
In [98]: # Compute the probability estimates for positive class
y_prob_svm = svm_model_linear.predict_proba(X_test_top20_svm)[: , 1]

# Compute fpr, tpr, and thresholds
fpr_svm, tpr_svm, thresholds_svm = roc_curve(y_test, y_prob_svm)

# Compute ROC AUC score
roc_auc_svm = roc_auc_score(y_test, y_prob_svm)
print('ROC AUC (SVM with linear kernel):', roc_auc_svm)

# Plot ROC AUC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr_svm, tpr_svm, color='blue', label='ROC Curve (AUC = {:.2f})'.format(roc_auc_svm))
plt.plot([0, 1], [0, 1], color='red', linestyle='--', label='Random')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - SVM with linear kernel (RFE-selected 20 features)')
plt.legend()
plt.show()
```

ROC AUC (SVM with linear kernel): 0.8258175398998815



EXPERIMENT 4

Using SMOTE Technique with RFE TOP 10 Features

Earlier we discovered that there is a class imbalance in the target variable we will be using the SMOTHE technique to address this class imabalance and compare results with other models done.

```
In [99]: # Apply SMOTE to handle class imbalance
smote = SMOTE(random_state=42)
X_train_smote, y_train_smote = smote.fit_resample(X_train_normalized, y_train)
```

Logistic Regression Model

```
In [100]: # Get Selected Features for each model using RFE with normalized data
logistic_regression_features = perform_rfe(logistic_regression, X_train_smote, y_train)
```

```
In [101]: # Train and evaluate Logistic Regression model with normalized data
logistic_regression.fit(X_train_smote[:, logistic_regression_features], y_train_smote)
y_pred_lr = logistic_regression.predict(X_test_normalized[:, logistic_regression_features])
```



```

accuracy_lr = accuracy_score(y_test, y_pred_lr)
conf_matrix_lr = confusion_matrix(y_test, y_pred_lr)

# Make predictions on the train set
y_pred_train_lr = logistic_regression.predict(X_train_smote[:, logistic_regression_fea

# Compute the accuracy of the train set for Logistic Regression model
train_accuracy_lr = accuracy_score(y_train_smote, y_pred_train_lr)
print("Logistic Regression Accuracy on the train set:", train_accuracy_lr)

# Print accuracy and classification report
print("Logistic Regression Accuracy:", accuracy_lr)
print(classification_report(y_test, y_pred_lr))

```

Logistic Regression Accuracy on the train set: 0.765859564164649

Logistic Regression Accuracy: 0.7391613361762616

	precision	recall	f1-score	support
0	0.90	0.72	0.80	1033
1	0.51	0.78	0.61	374
accuracy			0.74	1407
macro avg	0.70	0.75	0.71	1407
weighted avg	0.80	0.74	0.75	1407

In [102...


```

# Plot the confusion matrix
plt.figure()
cm_display_lr = ConfusionMatrixDisplay(confusion_matrix=conf_matrix_lr)
cm_display_lr.plot()
plt.title("Confusion Matrix - Logistic Regression with RFE(10) and SMOTE")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()

```

<Figure size 640x480 with 0 Axes>

Confusion Matrix - Logistic Regression with RFE(10) and SMOTE




In [103...

```
# Get predicted probabilities for positive class
y_prob_lr = logistic_regression.predict_proba(X_test_normalized[:, logistic_regression

# Compute false positive rate, true positive rate, and thresholds
fpr_lr, tpr_lr, thresholds_lr = roc_curve(y_test, y_prob_lr)

# Compute Area Under the Receiver Operating Characteristic Curve (ROC AUC)
auc_lr = roc_auc_score(y_test, y_prob_lr)

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr_lr, tpr_lr, label='Logistic Regression (area = %0.2f)' % auc_lr)
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve for Logistic Regression')
plt.legend(loc="lower right")
plt.show()
```



Random Forest Model

```
In [104... # Get Selected Features for Random Forest using RFE with normalized data
random_forest_features = perform_rfe(random_forest, X_train_smote, y_train_smote, n_fe

In [105... # Train Random Forest model with SMOTE data
random_forest.fit(X_train_smote[:, random_forest_features], y_train_smote)
y_pred_rf = random_forest.predict(X_test_normalized[:, random_forest_features])

# Evaluate Random Forest model
accuracy_rf = accuracy_score(y_test, y_pred_rf)
conf_matrix_rf = confusion_matrix(y_test, y_pred_rf)

# Make predictions on the train set
y_pred_train_rf = random_forest.predict(X_train_smote[:, random_forest_features])

# Compute the accuracy of the train set for Random Forest model
train_accuracy_rf = accuracy_score(y_train_smote, y_pred_train_rf)
print("Random Forest Accuracy on the train set with SMOTE data:", train_accuracy_rf)

print("Random Forest Accuracy:", accuracy_rf)
print("Random Forest Classification Report:")
print(classification_report(y_test, y_pred_rf))
```

Random Forest Accuracy on the train set with SMOTE data: 0.997457627118644

Random Forest Accuracy: 0.7484008528784648


Random Forest Classification Report:

	precision	recall	f1-score	support
0	0.84	0.81	0.82	1033
1	0.52	0.59	0.55	374
accuracy			0.75	1407
macro avg	0.68	0.70	0.69	1407
weighted avg	0.76	0.75	0.75	1407

In [106...

```
# Plot the Confusion Matrix for Random Forest Model
cm_display_rf = ConfusionMatrixDisplay(confusion_matrix=conf_matrix_rf, display_labels=
cm_display_rf.plot()
plt.title("Random Forest Confusion Matrix Using RFE Top 10 Features (Normalized Data)"
plt.show()
```

Random Forest Confusion Matrix Using RFE Top 10 Features (Normalized Data)




In [107...

```
# Plot the ROC curve for Random Forest Model
y_prob_rf = random_forest.predict_proba(X_test_normalized[:, random_forest_features]))[
fpr_rf, tpr_rf, thresholds_rf = roc_curve(y_test, y_prob_rf)
auc_rf = roc_auc_score(y_test, y_prob_rf)

plt.figure(figsize=(8, 6))
plt.plot(fpr_rf, tpr_rf, label='Random Forest (area = %0.2f)' % auc_rf)
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
```

```
plt.title('Receiver Operating Characteristic (ROC) Curve for Random Forest')
plt.legend(loc="lower right")
plt.show()
```



SVM Model

```
In [108... # Initialize RFE for feature selection with SVM model
rfe = RFE(estimator=svm_model_linear, n_features_to_select=10)

# Fit RFE on balanced data
rfe.fit(X_train_smote, y_train_smote)

# Get the selected features
selected_features_svm = rfe.support_
```

```
In [109... # Select top 10 features from training and test data
X_train_top10_svm = X_train_smote[:, selected_features_svm]
X_test_top10_svm = X_test_normalized[:, selected_features_svm]
```

```
In [110... # Train the SVM model on top 10 features
svm_model_linear.fit(X_train_top10_svm, y_train_smote)

# Make predictions
y_pred_svm = svm_model_linear.predict(X_test_top10_svm)
```

```

# Make predictions on the train set
y_pred_train_svm = svm_model_linear.predict(X_train_top10_svm)

# Compute the accuracy of the train set for SVM model
train_accuracy_svm = accuracy_score(y_train_smote, y_pred_train_svm)
print("SVM Accuracy on the train set with top 10 features and SMOTE data:", train_accu

# Evaluate the model
print("Scenario: SVM with RFE and SMOTE for dealing with class imbalance")
print("Classification Report:")
print(classification_report(y_test, y_pred_svm))
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred_svm))

```

SVM Accuracy on the train set with top 10 features and SMOTE data: 0.7456416464891041

Scenario: SVM with RFE and SMOTE for dealing with class imbalance

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.63	0.75	1033
1	0.45	0.83	0.58	374
accuracy			0.69	1407
macro avg	0.68	0.73	0.67	1407
weighted avg	0.79	0.69	0.70	1407

Confusion Matrix:

```

[[653 380]
 [ 63 311]]

```


In [111...

```

# Plot the confusion matrix for SVM
cm_display_svm = ConfusionMatrixDisplay(confusion_matrix=confusion_matrix(y_test, y_pr
cm_display_svm.plot()
plt.title("SVM Confusion Matrix Using RFE Top 10 Features and SMOTE (Normalized Data)"
plt.show()

```

SVM Confusion Matrix Using RFE Top 10 Features and SMOTE (Normalized Data)




In [112...

```
# Get the predicted probabilities for the positive class
y_prob_svm = svm_model_linear.predict_proba(X_test_top10_svm)[: , 1]

# Compute fpr, tpr, and thresholds
fpr, tpr, thresholds = roc_curve(y_test, y_prob_svm)

# Compute AUC score
auc = roc_auc_score(y_test, y_prob_svm)

# Plot ROC AUC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label="ROC Curve (AUC = {:.2f})".format(auc))
plt.plot([0, 1], [0, 1], linestyle="--", color="r", label="Random")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.legend()
plt.grid(False) # Remove gridlines
plt.show()
```



EXPERIMENT 5

Using SMOTE Technique with RFE TOP 20 Features

```
In [113... # Apply SMOTE to handle class imbalance
smote = SMOTE(random_state=42)
X_train_smote, y_train_smote = smote.fit_resample(X_train_normalized, y_train)
```

Logistic Regression Model

```
In [114... # Get Selected Features for each model using RFE with normalized data
logistic_regression_features = perform_rfe(logistic_regression, X_train_smote, y_train)
```

```
In [115... # Train and evaluate Logistic Regression model with normalized data
logistic_regression.fit(X_train_smote[:, logistic_regression_features], y_train_smote)
y_pred_lr = logistic_regression.predict(X_test_normalized[:, logistic_regression_features])
accuracy_lr = accuracy_score(y_test, y_pred_lr)
conf_matrix_lr = confusion_matrix(y_test, y_pred_lr)

# Make predictions on the train set
```



```

y_pred_train_lr = logistic_regression.predict(X_train_smote[:, logistic_regression_fea

# Compute the accuracy of the train set for Logistic Regression model
train_accuracy_lr = accuracy_score(y_train_smote, y_pred_train_lr)
print("Logistic Regression Accuracy on the train set with SMOTE data:", train_accuracy

# Print accuracy and classification report
print("Logistic Regression Accuracy:", accuracy_lr)
print(classification_report(y_test, y_pred_lr))

```

Logistic Regression Accuracy on the train set with SMOTE data: 0.7723970944309927
Logistic Regression Accuracy: 0.744136460554371

	precision	recall	f1-score	support
0	0.90	0.73	0.81	1033
1	0.51	0.78	0.62	374
accuracy			0.74	1407
macro avg	0.71	0.76	0.71	1407
weighted avg	0.80	0.74	0.76	1407


In [116...

```

# Plotting the Confusion Matrix for Logistic Regression with RFE and Smothe
cm_display_lr = ConfusionMatrixDisplay(conf_matrix_lr).plot()
plt.title("Logistic Regression Confusion Matrix Using SMOTHE and RFE (20)")
plt.show()

```

Logistic Regression Confusion Matrix Using SMOTHE and RFE (20)



In [117...

```

# Get predicted probabilities for positive class
y_prob_lr = logistic_regression.predict_proba(X_test_normalized[:, logistic_regression


# Compute false positive rate, true positive rate, and thresholds

```

```
fpr_lr, tpr_lr, thresholds_lr = roc_curve(y_test, y_prob_lr)

# Compute Area Under the Receiver Operating Characteristic Curve (ROC AUC)
auc_lr = roc_auc_score(y_test, y_prob_lr)

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr_lr, tpr_lr, label='Logistic Regression (area = %0.2f)' % auc_lr)
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve for Logistic Regression')
plt.legend(loc="lower right")
plt.show()
```



Random Forest Model

```
In [118... # Get Selected Features for Random Forest using RFE with normalized data
random_forest_features = perform_rfe(random_forest, X_train_smote, y_train_smote, n_fe

In [119... # Train Random Forest model with SMOTE data
random_forest.fit(X_train_smote[:, random_forest_features], y_train_smote)
y_pred_rf = random_forest.predict(X_test_normalized[:, random_forest_features])
```

```

# Make predictions on the train set
y_pred_train_rf = random_forest.predict(X_train_smote[:, random_forest_features])

# Compute the accuracy of the train set for Random Forest model
train_accuracy_rf = accuracy_score(y_train_smote, y_pred_train_rf)
print("Random Forest Accuracy on the train set with SMOTE data:", train_accuracy_rf)

# Evaluate Random Forest model
accuracy_rf = accuracy_score(y_test, y_pred_rf)
conf_matrix_rf = confusion_matrix(y_test, y_pred_rf)

print("Random Forest Accuracy:", accuracy_rf)
print("Random Forest Classification Report:")
print(classification_report(y_test, y_pred_rf))

```

Random Forest Accuracy on the train set with SMOTE data: 0.9984261501210654

Random Forest Accuracy: 0.7718550106609808

Random Forest Classification Report:

	precision	recall	f1-score	support
0	0.85	0.84	0.84	1033
1	0.57	0.57	0.57	374
accuracy			0.77	1407
macro avg	0.71	0.71	0.71	1407
weighted avg	0.77	0.77	0.77	1407


In [120...

```

# Plot the Confusion Matrix for Random Forest Model
cm_display_rf = ConfusionMatrixDisplay(confusion_matrix=conf_matrix_rf, display_labels=
cm_display_rf.plot()
plt.title("Random Forest Confusion Matrix Using RFE Top 10 Features (Normalized Data)")
plt.show()

```


Random Forest Confusion Matrix Using RFE Top 10 Features (Normalized Data)



In [121...

```
# Plot the ROC curve for Random Forest Model
y_prob_rf = random_forest.predict_proba(X_test_normalized[:, random_forest_features])[
fpr_rf, tpr_rf, thresholds_rf = roc_curve(y_test, y_prob_rf)
auc_rf = roc_auc_score(y_test, y_prob_rf)

plt.figure(figsize=(8, 6))
plt.plot(fpr_rf, tpr_rf, label='Random Forest (area = %0.2f)' % auc_rf)
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve for Random Forest')
plt.legend(loc="lower right")
plt.show()
```



SVM Model

```
In [122... # Initialize RFE for feature selection with SVM model
rfe = RFE(estimator=svm_model_linear, n_features_to_select=20)

# Fit RFE on normalized data
rfe.fit(X_train_smote, y_train_smote)

# Get the selected features
selected_features_svm = rfe.support_
```

```
In [123... # Select top 20 features from training and test data
X_train_top20_svm = X_train_smote[:, selected_features_svm]
X_test_top20_svm = X_test_normalized[:, selected_features_svm]
```

```
In [124... # Train the SVM model on top 20 features
svm_model_linear.fit(X_train_top20_svm, y_train_smote)

# Make predictions
y_pred_svm = svm_model_linear.predict(X_test_top20_svm)

# Make predictions on the train set
y_pred_train_svm = svm_model_linear.predict(X_train_top20_svm)
```

```

# Compute the accuracy of the train set for SVM model
train_accuracy_svm = accuracy_score(y_train_smote, y_pred_train_svm)
print("SVM Accuracy on the train set with top 20 features and SMOTE data:", train_accu

# Evaluate the model
print("SVM with RFE and SMOTE")
print("Classification Report:")
print(classification_report(y_test, y_pred_svm))
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred_svm))

```

SVM Accuracy on the train set with top 20 features and SMOTE data: 0.7452784503631962

SVM with RFE and SMOTE

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.63	0.75	1033
1	0.45	0.83	0.58	374
accuracy			0.69	1407
macro avg	0.68	0.73	0.67	1407
weighted avg	0.79	0.69	0.70	1407

Confusion Matrix:

```

[[653 380]
 [ 63 311]]

```


In [125...

```

# Plot the confusion matrix for SVM
cm_display_svm = ConfusionMatrixDisplay(confusion_matrix=confusion_matrix(y_test, y_pr
cm_display_svm.plot()
plt.title("SVM Confusion Matrix Using RFE Top 20 Features and SMOTE (Normalized Data)"
plt.show()

```

SVM Confusion Matrix Using RFE Top 20 Features and SMOTE (Normalized Data)




In [126...

```
# Get the predicted probabilities for the positive class
y_prob_svm = svm_model_linear.predict_proba(X_test_top20_svm)[: , 1]

# Compute fpr, tpr, and thresholds
fpr, tpr, thresholds = roc_curve(y_test, y_prob_svm)

# Compute AUC score
auc = roc_auc_score(y_test, y_prob_svm)

# Plot ROC AUC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label="ROC Curve (AUC = {:.2f})".format(auc))
plt.plot([0, 1], [0, 1], linestyle="--", color="r", label="Random")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("SVM ROC Curve With RFE Top 20 Features and SMOTE")
plt.legend()
plt.grid(False) # Remove gridlines
plt.show()
```



EXPERIMENT 6

Using SMOTE Technique On All Features

```
In [127... # Apply SMOTE to handle class imbalance
smote = SMOTE(random_state=42)
X_train_smote, y_train_smote = smote.fit_resample(X_train_normalized, y_train)
```

Logistic Regression Model

```
In [128... # Train and evaluate Logistic Regression model with SMOTE data
logistic_regression.fit(X_train_smote, y_train_smote)
y_pred_lr = logistic_regression.predict(X_test_normalized)
accuracy_lr = accuracy_score(y_test, y_pred_lr)
conf_matrix_lr = confusion_matrix(y_test, y_pred_lr)

# Make predictions on the train set
y_pred_train_lr = logistic_regression.predict(X_train_smote)

# Compute the accuracy of the train set for Logistic Regression model
train_accuracy_lr = accuracy_score(y_train_smote, y_pred_train_lr)
print("Logistic Regression Accuracy on the train set with SMOTE data:", train_accuracy)
```



```
# Print accuracy and classification report
print("Logistic Regression Accuracy:", accuracy_lr)
print("Logistic Regression Classification Report:")
print(classification_report(y_test, y_pred_lr))
```

Logistic Regression Accuracy on the train set with SMOTE data: 0.7730024213075061

Logistic Regression Accuracy: 0.7434257285003554


Logistic Regression Classification Report:

	precision	recall	f1-score	support
0	0.90	0.73	0.81	1033
1	0.51	0.78	0.62	374
accuracy			0.74	1407
macro avg	0.71	0.76	0.71	1407
weighted avg	0.80	0.74	0.76	1407

In [129...

```
# Plot Confusion Matrix
cm_display_lr = ConfusionMatrixDisplay(conf_matrix_lr).plot()
plt.title("Logistic Regression Confusion Matrix With SMOTHE On All Features")
plt.show()
```

Logistic Regression Confusion Matrix With SMOTHE On All Features



In [130...

```
# Calculate predicted probabilities
y_prob_lr = logistic_regression.predict_proba(X_test_normalized)[: , 1]

# Calculate ROC curve
fpr_lr, tpr_lr, thresholds_lr = roc_curve(y_test, y_prob_lr)

# Calculate AUC score
```


```

auc_lr = roc_auc_score(y_test, y_prob_lr)

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr_lr, tpr_lr, label='Logistic Regression (area = %0.2f)' % auc_lr)
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve With SMOTE on All Features for Logistic Regression')
plt.legend(loc="lower right")
plt.show()

```

Receiver Operating Characteristic (ROC) Curve With SMOTE on All Features for Logistic Regression



Random Forest Model

```

In [131... # Apply SMOTE to handle class imbalance
smote = SMOTE(random_state=42)
X_train_smote, y_train_smote = smote.fit_resample(X_train_normalized, y_train)

```

```

In [132... # Build and train Random Forest model
random_forest = RandomForestClassifier(random_state=42)
random_forest.fit(X_train_smote, y_train_smote)
# Predict on test set
y_pred_rf = random_forest.predict(X_test_normalized)

# Make predictions on the train set

```

```

y_pred_train_rf = random_forest.predict(X_train_smote)

# Compute the accuracy of the train set for Random Forest model
train_accuracy_rf = accuracy_score(y_train_smote, y_pred_train_rf)
print("Random Forest Accuracy on the train set with SMOTE data:", train_accuracy_rf)

# Evaluate the model
accuracy_rf = accuracy_score(y_test, y_pred_rf)
conf_matrix_rf = confusion_matrix(y_test, y_pred_rf)
print("Random Forest Accuracy:", accuracy_rf)
print("Random Forest Confusion Matrix:")
print(conf_matrix_rf)
print("Random Forest Classification Report:")
print(classification_report(y_test, y_pred_rf))

```

Random Forest Accuracy on the train set with SMOTE data: 0.9984261501210654

Random Forest Accuracy: 0.7626154939587776

Random Forest Confusion Matrix:

```
[[857 176]
```

```
[158 216]]
```

Random Forest Classification Report:

	precision	recall	f1-score	support
0	0.84	0.83	0.84	1033
1	0.55	0.58	0.56	374
accuracy			0.76	1407
macro avg	0.70	0.70	0.70	1407
weighted avg	0.77	0.76	0.76	1407

In [133...


```

# Plot Confusion Matrix
plt.figure(figsize=(8, 6))
cm_display_rf = ConfusionMatrixDisplay(confusion_matrix=conf_matrix_rf)
cm_display_rf.plot()
plt.title("Random Forest Confusion Matrix With SMOTE On All Features")
plt.show()

```

<Figure size 800x600 with 0 Axes>

Random Forest Confusion Matrix With SMOTE On All Features




In [134...

```
# Calculate AUC ROC for Random Forest
y_prob_rf = random_forest.predict_proba(X_test_normalized)[: , 1]
fpr_rf, tpr_rf, thresholds_rf = roc_curve(y_test, y_prob_rf)
auc_rf = roc_auc_score(y_test, y_prob_rf)

# Plot ROC curve for Random Forest
plt.figure(figsize=(8, 6))
plt.plot(fpr_rf, tpr_rf, label='Random Forest (area = %0.2f)' % auc_rf)
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve for Random Forest On All Features')
plt.legend(loc="lower right")
plt.show()
```

Receiver Operating Characteristic (ROC) Curve for Random Forest On All Features Using SMOTE



SVM Model

```
In [135... # Train the SVM model on all features
svm_model.fit(X_train_smote, y_train_smote)

# Make predictions
y_pred_svm = svm_model.predict(X_test_normalized)

# Make predictions on the train set
y_pred_train_svm = svm_model.predict(X_train_smote)

# Compute the accuracy of the train set for SVM model
train_accuracy_svm = accuracy_score(y_train_smote, y_pred_train_svm)
print("SVM Accuracy on the train set with SMOTE data:", train_accuracy_svm)

# Evaluate the model
print("SVM with SMOTE")
print("Classification Report:")
print(classification_report(y_test, y_pred_svm))
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred_svm))
```

SVM Accuracy on the train set with SMOTE data: 0.8504842615012107

SVM with SMOTE

Classification Report:

	precision	recall	f1-score	support
0	0.88	0.78	0.83	1033
1	0.54	0.72	0.61	374
accuracy			0.76	1407
macro avg	0.71	0.75	0.72	1407
weighted avg	0.79	0.76	0.77	1407

Confusion Matrix:


```
[[802 231]
```

```
[106 268]]
```

In [136...

```
# Plot the confusion matrix
plt.figure()
cm_display_svm = ConfusionMatrixDisplay(confusion_matrix=confusion_matrix(y_test, y_pr
cm_display_svm.plot()
plt.title("SVM Confusion Matrix Using All Features and SMOTE")
plt.show()
```

<Figure size 640x480 with 0 Axes>



In [137...

```
# Get the predicted probabilities for the positive class
y_prob_svm = svm_model.predict_proba(X_test_normalized)[: , 1]

# Compute fpr, tpr, and thresholds
fpr, tpr, thresholds = roc_curve(y_test, y_prob_svm)


# Compute AUC score
```

```

auc = roc_auc_score(y_test, y_prob_svm)

# Plot ROC AUC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label="ROC Curve (AUC = {:.2f})".format(auc))
plt.plot([0, 1], [0, 1], linestyle="--", color="r", label="Random")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - SVM on All Features with SMOTE")
plt.legend()
plt.show()

```



In the next step we will;

- Interpret the model outputs and draw conclusions based on the analysis.
- Communicate findings to stakeholders through reports, visualizations and presentations.
- Deployment of Models
- Outline actionable insights derived from the analysis into decision-making processes.